
connectal Documentation

Release 16.05.2

Jamey Hicks, Myron King, John Ankcorn

August 18, 2016

CONTENTS

1	Contents	1
1.1	Introduction	1
1.2	Installation	1
1.3	Connectal Design	2
1.4	Connectal Developer's Guide	29
1.5	Connectal Makefile Variables	34
1.6	Connectal Tools	38
1.7	Connectal BSV Libraries	45
1.8	Connectal C/C++ Libraries	60
1.9	Connectal Examples	62
1.10	Indices and tables	62
	Bsv Package Index	63
	Index	65

CONTENTS

Introduction

Introduction goes here.

Installation

Installing Connectal from Packages

On Ubuntu systems, connectal may be installed from pre-built packages:

```
sudo add-apt-repository -y ppa:jamey-hicks/connectal sudo apt-get update sudo apt-get -y install connectal
```

Installing Connectal from Source

Connectal source comes from three repositories:

```
git clone git://github.com/cambridgehackers/connectal git clone git://github.com/cambridgehackers/fpgamake  
git clone git://github.com/cambridgehackers/buildcache
```

To use Connectal to build hardware/software applications, some additional packages are required:

```
cd connectal; sudo make install-dependences
```

Installing Connectal Drivers and PCI Express Utilities:

To run Connectal applications on FPGAs attached via PCI Express, a couple of device drivers have to be built and installed:

```
cd connectal; make all && sudo make install
```

In addition, you will need to build and install fpgajtag and pciescan:

```
git clone git://github.com/cambridgehackers/pciescan (cd pciescan; make && sudo make install)
```

```
git clone git://github.com/cambridgehackers/fpgajtag (cd fpgajtag; make && sudo make install)
```

Installing Vivado

Download Vivado from Xilinx

<http://www.xilinx.com/support/download/index.html/content/xilinx/en/downloadNav/vivado-design-tools/2015-4.html>

Connectal builds do not use the Vivado SDK.

Connectal Design

Abstract

The cost and complexity of hardware-centric systems can often be reduced by using software to perform tasks which don't appear on the critical path. Alternately, the performance of software can sometimes be improved by using special purpose hardware to implement tasks which *do* appear on the critical path. Whatever the motivation, most modern systems are composed of both hardware and software components.

Given the importance of the connection between hardware and software in these systems, it is surprising how little automated and machine-checkable support there is for co-design space exploration. This paper presents the Connectal framework, which enables the development of hardware accelerators for software applications by generating hardware/software interface implementations from abstract Interface Design Language (IDL) specifications.

Connectal generates stubs to support asynchronous remote method invocation from software to software, hardware to software, software to hardware, and hardware to hardware. For high-bandwidth communication, the Connectal framework provides comprehensive support for shared memory between hardware and software components, removing the repetitive work of processor bus interfacing from project tasks.

This framework is released as open software under an MIT license, making it available for use in any projects.

Introduction

Because they are so small and inexpensive, processors are now included in all but the smallest hardware designs. This grants flexibility to hardware designers because the non-performance-critical components can be implemented in software and the performance-critical components can be implemented in hardware. Using software for parts of the design can decrease the effort required to implement configuration and orchestration logic (for example). It can also offer hardware developers greater adaptability in meeting new project requirements or supporting additional applications.

As a system evolves through design exploration, the boundary between the software and hardware pieces can change substantially. The old paradigm of “separate hardware and software designs before the project starts” is no longer sustainable, and hardware teams are increasingly responsible for delivering significant software components.

Despite this trend, hardware engineers find themselves with surprisingly poor support for the development of the software that is so integral to their project's success. They are often required to manually develop the necessary software and hardware to connect the two environments. In the software world, this is equivalent to manually re-creating header files from the prose description of an interface implemented by a library. Such ad hoc solutions are tedious, fragile, and difficult to maintain. Without a consistent framework and toolchain for jointly managing the components of the hardware/software boundary, designers are prone to make simple errors which can be expensive to debug.

The goal of our work is to support the flexible and consistent partitioning of designs across hardware and software components. We have identified the following four goals as central to this endeavor:

- Connect software and hardware by compiling interface declarations.

- Enable concurrent access to hardware accelerators from software.
- Enable high-bandwidth sharing of system memory with hardware accelerators.
- Provide portability across platforms (CPU, OS, bus types, FPGAs).

In this paper, we present a software-driven hardware development framework called [Connectal](#). Connectal consists of a fully-scripted tool-chain and a collection of libraries which can be used to develop production quality applications comprised of software components running on CPUs communicating with hardware components implemented in FPGA or ASIC.

When designing Connectal, our primary goal was to create a collection of components which are easy to use for simple implementations and which can be configured or tuned for high performance in more complicated applications. To this end, we adopted a decidedly minimalist approach, attempting to provide the smallest viable programming interface which can guarantee consistent access to shared resources in a wide range of software and hardware execution environments. Because our framework targets the implementation of performance-critical systems rather than their simulation, we have worked hard to remove any performance penalty associated with its use.

We wrote the hardware components of the Connectal libraries in [Bluespec System Verilog](#) (BSV) because it enables a higher level of abstraction than the alternatives and supports parameterized types. The software components are implemented in C/C++. We chose Bluespec interfaces as the interface definition language (IDL) for Connectal’s interface compiler.

This paper describes the Connectal framework, and how it can be used to flexibly move between a variety of software environments and communication models when mapping applications to platforms with connected FPGAs and CPUs.

Document Organization

In Section [Accelerating String Search](#), we present an example running in a number of different execution environments. In Section [The Connectal Framework](#), we give an overview of the Connectal framework and its design goals. In Section [Sec-Impl](#) we discuss the details of Connectal and how it can be used to implement the example. Section [Sec-ToolChain](#) describes the implementation of Connectal, supported platforms, and the tool chain used to coordinate the various parts of the framework. The paper concludes with a discussion of performance metrics and related work.

Accelerating String Search

The structure of a hardware/software (HW/SW) system can evolve quite dramatically to reflect changing requirements, or during design exploration. In this section, we consider several different implementations of a simple string search application~cite{mpAlgo}. Each variation represents a step in the iterative refinement process, intended to enhance performance or enable a different software execution environment.

```

begin{figure}[!h]
    centering
    includegraphics[width=0.43textwidth]{platform.pdf}
    caption{label{Fig:Platform0}Target platform for string search application}
end{figure}

```

Figure~ref{Fig:Platform0} shows the target platform for our example. The pertinent components of the host system are the multi-core CPU, system memory, and PCI Express (PCIe) bus. The software components of our application will be run on the CPU in a Linux environment. Connected to the host is a PCIe expansion card containing (among other things) a high-performance FPGA chip and a large array of flash memory. The FPGA board was designed as a platform to accelerate “big data” analytics by moving more processing power closer to the storage device.

Initial Implementation

```
begin{figure}[!h]
    centering
    includegraphics[width=0.43textwidth]{data_accel_logical0.pdf}
    caption{label{Fig:StringSearch0}Logical components of the string search system}
end{figure}
```

The design process really begins with a pure software implementation of the algorithm, but the first attempt we consider is the initial inclusion of HW acceleration shown in Figure~ref{Fig:StringSearch0}. The search functionality is executed by software running in user-space which communicates with the hardware accelerator through a device driver running in the Linux kernel. The hardware accelerator, implemented in the FPGA fabric, executes searches over data stored in the flash array as directed by the software.

The FPGA has direct access to the massive flash memory array, so if we implement the search kernel in hardware, we can avoid bringing data into the CPU cache (an important consideration if we intend to run other programs simultaneously). By exploiting the high parallelism of the execution fabric as well as application aware caching of data, an FPGA implementation can outperform the same search executed on the CPU.

Multithreading the Software

The efficient use of flash memory requires a relatively sophisticated management strategy. Our first refinement is based on the observation that there are four distinct tasks which the application software executes (mostly) independently:

- Send search command to the hardware.
- Receive search results from the hardware.
- Send commands to the hardware to manage the flash arrays
- Receive responses from the flash management hardware

To exploit the task-level parallelism in our application, we can assign one thread to each of the four enumerated tasks. To further improve efficiency, the two threads receiving data from the hardware put themselves to sleep by calling `textbf{poll}` and are woken up only when a message has been received.

```
begin{figure}[!h]
    centering
    includegraphics[width=0.43textwidth]{data_accel_logical1.pdf}
    caption{label{Fig:StringSearch1}Using a mutex to coordinate user-level access to hardware accelerator}
end{figure}
```

With the introduction of multithreading, we will need a synchronization mechanism to enforce coherent access to the hardware resources. Because the tasks which need coordinating are all being executed as user-space threads, the access control must be implemented in software as well. As shown in Figure~ref{Fig:StringSearch1}, a mutex is used to coordinate access to the shared hardware resource between user-level processes.

Refining the Interfaces

```
begin{figure}[!h] centering includegraphics[width=0.43textwidth]{data_accel_logical2.pdf} caption{label{Fig:StringSearch2}Movement of functionality from
```


user to kernel space. Software-based coordination between kernel and user processes are prohibitively expensive.}

end{figure}

Figure~ref{Fig:StringSearch2} shows a further refinement to our system in which we have reimplemented the Flash Management functionality as a block-device driver. Instead of directly operating on physical addresses, the string search now takes a file descriptor as input and uses a Linux system-call to retrieve the file block addresses through the file system. This refinement permits other developers to write applications which can take advantage of the accelerator without any knowledge of the internal details of the underlying storage device. It also enables support for different file systems as we now use a POSIX interface to generate physical block lists for the the storage device hardware. The problem with this refinement is that we no longer have an efficient SW mechanism to synchronize the block device driver running in kernel space with the application running in user space.

begin{figure}[htb] centering includegraphics[width=0.43textwidth]{data_accel_logical3.pdf} caption{label{Fig:StringSearch3}Correct interface design

removes the need for coordination between user and kernel threads.}

end{figure}

To solve to this problem (shown in Figure~ref{Fig:StringSearch3}), we can remove the need for explicit SW coordination altogether by giving each thread uncontested access to its own dedicated HW resources mapped into disjoint address regions. (There will of course be implicit synchronization through the file system.)

Shared Access to Host Memory

In the previous implementations, all communication between hardware and software takes place through memory mapped register IO. Suppose that instead of searching for single strings, we want to search for large numbers of (potentially lengthy) strings stored in the flash array. Attempting to transfer these strings to the hardware accelerator using programmed register transfers introduces a performance bottleneck. In our final refinement, the program will allocate memory on the host system, populate it with the search strings, and pass a reference to this memory to the hardware accelerator which can then read the search strings directly from the host memory.

begin{figure}[htb] centering includegraphics[width=0.43textwidth]{data_accel_logical4.pdf} caption{label{Fig:StringSearch4}Connectal support for DMA.}

end{figure}

Efficient high-bandwidth communication in this style requires the ability to share allocated memory regions between hardware and software processes without copying. Normally, a programmer would simply call application space `textbf{malloc}`, but this does not provide a buffer that can be shared with hardware or other software processes. As shown in Figure~ref{Fig:StringSearch4}, a special-purpose memory allocator has been implemented in Linux, using `dmabufciter{dmabuf}` to provide reference counted sharing of memory buffers across user processes and hardware.

To conclude, we consider how the HW/SW interface changed to accommodate each step in the refinement process: The hardware interface required by the design in Figure~ref{Fig:StringSearch0} is relatively simple. Command/response queues in the hardware accelerator are exposed using a register interface with accompanying *empty/full* signals. To support the use of *poll* by the refinement in Figure~ref{Fig:StringSearch1}, interrupt signals must be added to the hardware interface and connected to the Linux kernel. Partitioning the address space as required by the refinement in Figure~ref{Fig:StringSearch3} necessitates a consistent remapping of registers in both hardware and software.

The Connectal Framework

In and of themselves, none of the HW/SW interfaces considered in Section *Accelerating String Search* are particularly complex. On the other hand, implementing the complete set and maintaining correctness as the application evolves is a considerable amount of care, requiring deep understanding of both the application and the platform. The Connectal

framework is a collection of tools and library components which was designed to address these challenges with the following features:

- Easy declaration and invocation of remote methods between application components running on the host or in the FPGA.
- Direct user-mode access to hardware accelerators from software.
- High performance read and write bus master access to system memory from the FPGA
- Infrastructure for sharing full speed memory port access between an arbitrary number of clients in the FPGA fabric
- Portability across platforms using different CPUs, buses, operating systems, and FPGAs
- Fully integrated tool-chain support for dependency builds and device configuration.

In this section, we introduce the Connectal framework through a discussion of its prominent features.

Portals

Connectal implements remote method invocation between application components using asynchronous messaging. The message and channel types are application specific, requiring the user to define the HW/SW interface using BSV interfaces as the interface definition language (IDL). These interfaces declare logical groups of unidirectional “send” methods, each of which is implemented as a FIFO channel by the Connectal interface compiler; all channels corresponding to a single BSV interface are grouped together into a single *portal*.

From the interface specification, the Connectal interface compiler generates code for marshalling the arguments of a method into a message to be sent and unmarshaling values from a received message. It generates a `textit{proxy}` to be invoked on the sending side and a `textit{wrapper}` that invokes the appropriate method on the receiving side. Platform specific libraries are used to connect the proxies and wrappers to the communication fabric.

In the hardware, each portal is assigned a disjoint address range. On the host, Connectal assigns each portal a unique Linux device (`/dev/portal*n*$`) which is accessed by the application software using the generated wrappers and proxies. An application can partition methods across several portals, to control access to the interfaces by specific hardware or software modules. To support bi-directional communication, at least two portals are required: one which allows software to “invoke” hardware, and another for hardware to “invoke” software. Each portal may be accessed by different threads, processes, or directly from the kernel.

Direct user-mode access to hardware

We designed Connectal to provide direct access to accelerators from user-mode programs in order to eliminate the need for device-drivers specific to each accelerator. We have implemented a kernel module for both X86 and ARM architectures with a minimal set of functionality: the driver implements `textbf{mmap}` to map hardware registers into user space and `textbf{poll}` to enable applications to suspend a thread waiting for interrupts originating from the hardware accelerators. These two pieces of functionality have been defined to be completely generic; no modification is required to kernel drivers as the HW/SW interface evolves. All knowledge of the interface register semantics (and corresponding changes) is encoded by the interface compiler in the generated proxies and wrappers which are compiled as part of the application and executed in user-mode.

This approach is known as user-space device drivers~cite{Khalidi:1995:EZI:974947,UIO:Howto} and has a number of distinct advantages over traditional kernel modules. To begin with, it reduces the number of components that need to be modified if the HW/SW interface changes, and eliminates the need for device-driver development expertise in many cases. Secondly, after the hardware registers have been mapped into user address space, the need for software to switch between user and kernel mode is all but eliminated since all “driver” functionality is being executed in user-space.

Shared Access to Host Memory

Connectal generates a hardware FIFO corresponding to each method in the portal interface, and the software reads and writes these FIFOs under certain conditions. To improve throughput, Connectal libraries also support credit-based flow-control. Though credit-based flow-control with interrupts is more efficient than polling status registers from software, there is often the need for much higher bandwidth communication between the hardware and software.

Hardware accelerators often communicate with the application through direct access to shared memory. An important feature of Connectal is a flexible, high performance API for allocating and sharing such memory, and support for reading and writing this memory from hardware and software. The Connectal framework implements this through the combination of a Linux kernel driver, C++ libraries, and BSV modules for the FPGA. We implemented a custom kernel memory allocator for Connectal, `textbf{portalmem}`, using the kernel `dmabuf` support. Any solution which allocates and shares memory between hardware and software must meet two high-level requirements:

- **Allocated buffers must have reference counts to prevent memory leaks.**
- **Efficient mechanisms must be provided to share the location of** allocated regions.

Using the `portalmem` driver, programs can allocate regions of system memory (DRAM) and map it into their own virtual address space. Reference-counted access to shared memory regions allocated using `portalmem` can be granted to other SW processes by transmitting the file descriptor for the allocated region. Reference counting has been implemented in the driver so that once an allocated memory region has been dereferenced by all SW and HW processes, it will be deallocated and returned to the kernel free memory pool.

Simple hardware accelerators often require contiguous physical addresses. Unfortunately, when allocating memory from a shared pool in a running system, obtaining large areas of contiguous memory is often problematic, limiting the size of the region that can be allocated. To support indexed access to non-contiguous memory aggregates, Connectal provides address translation support to hardware accelerators in the FPGA, similar to the MMU functionality on the CPU side.

Distributed Access to Memory Ports

When building accelerators for an algorithm, multiple parameters are often accessed directly from system memory using DMA. As the hardware implementation is parallelized, multiple accesses to each parameter may be required. In these cases, the number of memory clients in the application hardware usually exceeds the number of host memory ports. Sharing these ports requires substantial effort, and scaling up a memory interconnect while maximizing throughput and clock speed is extremely challenging.


To support this common design pattern, the Connectal framework provides a portable, scalable, high performance library that applications can use to facilitate the efficient sharing of host memory ports. This library is implemented as parameterized Bluespec modules which allow the user to easily configure high-performance memory access trees, supporting both reading and writing.

Platform Portability

We structured Connectal to improve the portability of applications across CPU types, operating systems, FPGAs, and how the CPU and FPGA are connected. The software and hardware libraries are largely platform independent. As a result, applications implemented in the framework can be compiled to run on the range of different platforms.

Supported platforms are shown in Figure Fig-platforms. Application software can be executed on x86 and ARM CPUs running either Ubuntu or Android operating systems. A range of different Xilinx FPGAs can be connected to the CPU and system memory via PCI Express or AXI. The BSV simulator (Bluesim) can be used in place of actual FPGA hardware for debugging purposes.

When the target application needs to interact with other Linux kernel resources (for example, a block device or a network interface), the application may run in kernel mode with the logic run either in an FPGA or in Bluesim.



Implementing String Search

Having covered the features of the Connectal at a high level, we now explain more specifically how the framework can be applied to implement the refinements outlined in Section *Accelerating String Search*.

Initial Implementation

The FPGA is connected to the host system with a PCIe bus, and to the memory array with wires. In addition to implementing a search kernel, the hardware accelerator must communicate with the software components and with the flash chips. Communication with the software takes place through portals, whose interface declaration is given below:

```

interface StrstrRequest;
    method Action setupNeedle(Bit#(8) needleChars);
    method Action search(Bit#(32) haystackPtr,
                        Bit#(32) haystackLen);
endinterface
interface StrstrIndication;
    method Action searchResult(Int#(32) v);
    method Action setupComplete();
endinterface

```

The hardware implements the `StrstrRequest` interface, which the software invokes (remotely) to specify the search string and the location in flash memory to search. The software implements the `StrstrIndication` interface, which the hardware invokes (remotely) to notify the software of configuration completion or search results. The interface compiler generates a separate portal for each of these interfaces. Within each portal, a dedicated unidirectional FIFO is assigned to each logical interface method.

In our initial implementation the accelerator does not access system memory directly, so the search string is transmitted to the accelerator one character at a time via the `{tt setupNeedle}` method. We will see in Section :ref:Sec-StringSearchSystemMemory how to use a pointer to system memory instead.

Invoking Hardware from Software

Because the `StrstrRequest` functionality is implemented in hardware, the Connectal interface compiler generates a C++ `textbf{proxy}` with the following interface to be invoked by the application software:

```

class StrstrRequestProxy : public Portal {
public:
    void setupNeedle(uint32_t needleChars);
    void search(uint32_t haystackPtr,
               uint32_t haystackLen);
};

```

The implementation of `StrStrRequestProxy` marshals the arguments of each method and en-queues them directly into their dedicated hardware FIFOs. To execute searches in the FPGA fabric over data stored in flash memory, the software developer simply instantiates `StrStrRequestProxy` and invokes its methods:

```
StrStrRequestProxy *proxy =
    new StrStrRequestProxy(...);
proxy->search(haystackPtr, haystackLen);
```

On the FPGA, the user implements the application logic as a BSV module with the `StrStrRequest` interface. A *wrapper* is generated by the interface compiler to connect this module to the hardware FIFOs. The wrapper unmarshals messages that it receives and then invokes the appropriate method in the `StrStrRequest` interface. Here is the BSV code that instantiates the generated wrapper and connects it to the user's `texttt{mkStrStr}` module:

```
StrStrRequest strStr <- mkStrStr(...);
StrStrRequestWrapper wrapper <-
    mkStrStrRequestWrapper(strStr);
```

Figure [:ref:'Fig-msc1'](#) shows how all the pieces of an application implemented using Connectal work together when hardware functionality is invoked remotely from software. Direct access to the memory mapped hardware FIFOs by the generated proxy running in user-mode is key to the efficiency of our implementation strategy.

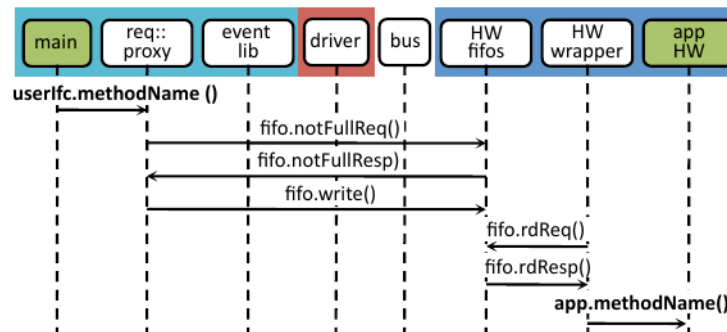


Fig. 1.1: SW invokes HW: *main* and *app HW* are implemented by the user.

Invoking Software from Hardware

Invoking software from hardware takes a slightly different form, due primarily to the fact that “main” is still owned by software. Since the direction of the remote invocation is reversed, the proxy on this path will be instantiated on the FPGA and the wrapper instantiated on host side. The user implements the `StrStrResponse` interface in software and connects it to the generated wrapper using C++ subclasses:

```
class StrStrResponse:
public StrStrResponseWrapper {
    ...
    void searchResult(int32_t v) {...}
}
```

The `StrStrResponseWrapper` constructor registers a pointer to the object with the event library which keeps track of all instantiated software wrappers. The wrapper implementation unmarshals messages sent through the hardware FIFOs and invokes the appropriate subclass interface method. To activate this path, main simply instantiates the response implementation and invokes the library event handler:

```
StrStrResponse *response =
    new StrStrResponse(...);
```

```
while(1)
  portalExec_event();
```

On the invocation side, the interface compiler generates a proxy which the application logic instantiates and invokes directly:

```
StrStrResponseProxy proxy <-
  mkStrStrRequestProxy();
StrStrRequest strStr <-
  mkStrStr(... proxy.ifc ...);
```

Figure :ref:Fig-msc0 shows how all the pieces of an application collaborate when software functionality is being invoked from hardware.

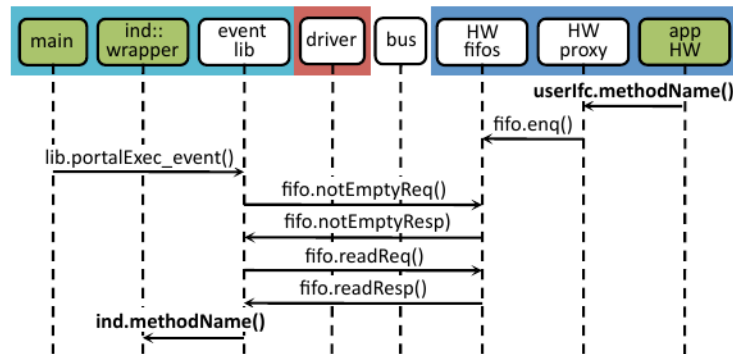


Fig. 1.2: HW invokes SW: ‘main’, ‘ind::wrapper’, and ‘app HW’ are implemented by the user.

The simplest software execution environment for the string search accelerator is to have a single thread making requests and waiting for responses as follows:

```
void search(char *str){
  StrStrRequestProxy *req =
    new StrStrRequestProxy(...);
  StrStrResponse *resp =
    new StrStrResponse(...);
  while (char c = *str++)
    req->setupNeedle(c);
  // start search
  req->search(...);
  // handle responses from the HW
  while(1)
    portalExec_event();
}
```

The call to `:c:func:portalExec_event()` checks for a response from HW. If there is a pending response, it invokes the method corresponding to that FIFO in the wrapper class. This generated method reads out a complete message from the FIFO and unmarshals it before invoking the user-defined call-back function, which in this case would be `texttt{StrStrResponse::searchResult}`.

Connecting To Flash

On BlueDBM, one of our target platforms, the flash memory array is connected directly to the FPGA chip, and DDR signals are used to read/write/erase flash memory cells. The RTL required to communicate with the memory requires

some commonly used functionality, such as *SerDes* and DDR controllers, both of which are included in the BSV libraries distributed as part of the Connectal framework.

Multithreading The Software

In many cases, we would like to avoid a hardware-to-software path which requires the software to poll a hardware register on the other side of a bus for relatively infrequent events. To accommodate this, the Connectal framework generates interrupts which are raised when hardware invokes software interface methods. The generic Connectal driver connects these signals to the Linux kernel and the software wrappers can exploit then by calling poll. Connectal applications often use a separate thread to execute hardware-to-software asynchronous invocations, since dedicated thread can put itself to sleep until the hardware raises an interrupt. The “main” thread is free to do other work and can communicate with the “indication” thread using a semaphore as shown below:

```
class StrStrResponse:
    public StrStrResponseWrapper {
    sem_t sem;
    int v;
    void searchResult(int32_t v) {
        this->response = v;
        sem_post(&sem);
    }
    void waitResponse() {sem_wait(&sem);}
};
StrStrResponse *resp;
StrStrRequestProxy *req;
int search(char *str){
    while (char c = *str++)
        req->setupNeedle(c);
    // start search
    req->search(...);
    // wait for response
    resp->waitResponse();
    // return result
    return resp->v;
}
```

The polling thread is started by a call to `:c:func:portalExec_start()`, which ultimately invokes the `:c:func:portalExec_poll()` function implemented in the Connectal event library. `:c:func:portalExec_poll()` invokes the system call `textbf{poll}` on the FDs corresponding to all the indication or response portals, putting itself to sleep. When an interface method is invoked in the hardware proxy, an interrupt is raised, waking the indication thread. A register is read which indicates which method is being called and the corresponding wrapper method is invoked to read/marshal the arguments and invoke the actual user-defined methods. Figure Fig-msc2 shows this process. Multithreading often leads to simultaneous access to shared hardware resources. If a software solution to protect these resources (such as mutex) is not available, the hardware interface can be refactored into separate portals, one for each control thread.

Each interface will generate a separate Portal which is assigned its own address space and Linux device. Using Linux devices in this way enables access control restrictions to be specified individually for each portal. This feature can be used to grant different users or processes exclusive access and prevent unauthorized access to specific pieces of hardware functionality.

Shared Access to Host Memory

In the first three refinements presented in Section [Accelerating String Search](#), all communication between hardware and software takes place through register-mapped IO. The final refinement in Section `Sec-StrStrDma` is to grant hardware and software shared access to host memory. The interface to the search accelerator shown below has been updated to use direct access to system memory for the search strings:


```
interface StrstrRequest;
  method Action setup(Bit#(32) needlePtr,
                    Bit#(32) mpNextPtr,
                    Bit#(32) needleLen);
  method Action search(Bit#(32) haystackPtr,
                    Bit#(32) haystackLen,
                    Bit#(32) iterCount);
endinterface
interface StrstrIndication;
  method Action searchResult(Int#(32) v);
  method Action setupComplete();
endinterface
```

In order to share memory with hardware accelerators, it needs to be allocated using `:c:func:portalAlloc()`. Here is the search function updated accordingly:

```
int search(char *str){
  int size = strlen(str)+1;
  int fd = portalAlloc(size, 0);
  char *sharedStr = portalMmap(fd, size);
  strcpy(sharedStr, str);
  // send a DMA reference to the search pattern
  req->needle(dma->reference(fd), size);
  // start search
  req->search(...);
  resp->waitResponse();
  ... unmap and free the string
  return resp->v;
}
```

The application allocates shared memory via `{tt portalAlloc}`, which returns a file descriptor, and then passes that file descriptor to `{tt mmap}`, which maps the physical pages into the application's address space. The file descriptor corresponds to a `dmabufcite{dmabuf}`, which is a standard Linux kernel mechanism.

To share that memory with the accelerator, the application calls `{tt reference}`, which sends a logical to physical address

mapping to the hardware's address translator. The call to `{tt reference}` returns a handle, which the application sends to

the accelerator. Connectal's BSV libraries for DMA enable the accelerator to read or write from offsets to these handles, taking care of address translation transparently.

To fully exploit the data parallelism, `{tt mkStrStr}` partitions the search space into `p` partitions. It instantiates two memory read trees from the Connectal library (`{tt MemReadEngineV}`), discussed in Section [Distributed Access to Memory Ports](#), each with `p` read servers. One set is used by the search kernels to read the configuration data from the host memory, while the other is used to read the “haystack” from flash.

On supported platforms such as Zynq which provide multiple physical master connections to system memory, Connectal interleaves DMA requests over the parallel links. It does this on a per-read-client basis, rather than a per-request basis.

Alternate Portal Implementations

Connectal separates the generation of code for marshalling and unmarshalling method arguments from the transport mechanism used to transmit the messages. This separation enables “swappable” application-specific transport libraries. In light of this, a large number of transport mechanism can be considered. Switching between mechanism requires a simple directive in the project Makefile (more details are given in Section [Sec-ToolChain](#)).

By default, each portal is mapped to a region of address space and a memory-mapped FIFO channel is generated for each method. Though software access to all FIFO channels in a design may occur through single bus slave interface, Connectal libraries implement their multiplexing to ensure that each FIFO is independent, allowing concurrent access to different methods from multiple threads or processes.

The default portal library implements the method FIFOs in the hardware accelerator. This provides the lowest latency path between hardware and software, taking about 1 microsecond to send a message. If higher bandwidth or transaction rates are needed, FIFOs implemented as a ring buffer in DRAM can be used instead. This requires more instructions per message send and receive, but may achieve higher throughput between the CPU and hardware.

During the design exploration process, a component originally implemented on the FPGA may migrate to software running on the host processor. Remote invocations which were originally from software to hardware must be recast as software to software. Without changing the IDL specification, the transport mechanism assigned to a portal can be re-specified to implement communication between software components running either on the same host or across a network.

Connectal uses UNIX sockets or shared memory to transport messages between the application software components or the hardware simulator. In other situations, TCP or UDP can be used to transport the messages to hardware running on another machine. Viable connections to the FPGA board range from low-speed interconnects such as JTAG, SPI, to higher-speed interconnects such as USB or Aurora over multi-gigabit per second transceivers.

Workflow using Connectal

In this section, we give an overview of the Connectal workflow and toolchain. The complete toolchain, libraries, and many running examples may be obtained at [textit{www.connectal.org}](http://www.connectal.org) or by emailing tex-tit@googlegroups.com.

Top level structure of Connectal applications

The simplest Connectal application consists of 4 files:

Makefile

The top-level Makefile defines parameters for the entire application build process. In its simplest form, it specifies which Bluespec interfaces to use as portals, the hardware and software source files, and the libraries to use for the hardware and software compilation.

Application Hardware

Connectal applications typically have at least one BSV file containing declarations of the interfaces being exposed as portals, along with the implementation of the application hardware itself.

Top.bsv

In this file, the developer instantiates the application hardware modules, connecting them to the generated wrappers and proxies for the portals exported to software. To connect to the host processor bus, a parameterized standard interface is used, making it easy to synthesize the application for different CPUs or for simulation. If CPU specific interface signals are needed by the design (for example, extra clocks that are generated by the PCIe core), then an optional CPU-specific interface can also be used.

If the design uses multiple clock domains or additional pins on the FPGA, those connections are also made here by exporting a ‘Pins’ interface. The Bluespec compiler generates a Verilog module from the top level BSV module, in

which the methods of exposed interfaces are implemented as Verilog ports. Those ports are associated to physical pins on the FPGA using a physical constraints file.

Application CPP

The software portion of a Connectal application generally consists of at least one C++ file, which instantiates the generated software portal wrapper and proxies. The application software is also responsible for implementing main.

Development cycle

After creating or editing the source code for the application, the development cycle consists of four steps: generating makefiles, compiling the interface, building the application, and running the application.

Generating Makefiles

Given the parameters specified in the application Makefile and a platform target specified at the command line, Connectal generates a target-specific Makefile to control the build process. This Makefile contains the complete dependency information for the generation of wrappers/proxies, the use of these wrappers/proxies in compiling both the software and hardware, and the collection of build artifacts into a package that can be either run locally or over a network to a remote ‘device under test’ machine.

Compiling the Interface

The Connectal interface compiler generates the C++ and BSV files to implement wrappers and proxies for all interfaces specified in the application Makefile. Human readable textbf{JSON} is used as an intermediate representation of portal interfaces, exposing a useful debugging window as well as a path for future support of additional languages and IDLs.

Building the Application

A target in the generated Makefile invokes GCC to compile the software components of the application. The Bluespec compiler (bsc) is then invoked to compile the hardware components to Verilog. A parameterized Tcl script is used to drive Vivado to build the Xilinx FPGA configuration bitstream for the design.

A Connectal utility called `fpgamake` supports specification of which Bluespec and Verilog modules should be compiled to separate netlists and to enable separate place and route of those netlists given a floor plan. Separate synthesis and floor planning in this manner can reduce build times, and to make it easier to meet timing constraints.

Another Connectal utility called `buildcache` speeds recompilation by caching previous compilation results and detecting cases where input files have not changed. Although similar to the better-known utility `textit{ccache}`, this program has no specific knowledge of the tools being executed, allowing it to be integrated into any workflow and any tool set. This utility uses the system call `textbf{strace}` to track which files are read and written by each build step, computing an ‘input signature’ of the MD5 checksum for each of these files. When the input signature matches, the output files are just refreshed from the cache, avoiding the long synthesis times for the unchanged portions of the project.

Running the Application

As part of our goal to have a fully scripted design flow, the generated Makefile includes a `texttt{run}` target that will program the FPGA and launch the specified application or test bench. In order to support shared target hardware resources, the developer can direct the run to a particular machine, which can be accessed over the network. For Ubuntu target machines, `ssh` is used to copy/run the application. For Android target machines, ‘adb’ is used.

Continuous Integration and Debug Support

Connectal provides a fully scripted flow in order to make it easy to automate the building and running of applications for continuous integration. Our development team builds and runs large collections of tests whenever the source code repository is updated.

Connectal also provides trace ring buffers in hardware and analysis software to trace and display the last transactions on the PCIe or AXI memory bus. This trace is useful when debugging performance or correctness problems, answering questions of the form:

- What were the last memory requests and responses?
- What was the timing of the last memory request and responses?
- What were the last hardware method invocations or indications?

Performance of Generated Systems

A framework is only useful if it reduces the effort required by developers to achieve the desired performance objective. Trying to gauge the relative effort is difficult since the authors implemented both the framework and the running example. On PCIe-based platforms we were able to reduce the time required to search for a fixed set of strings in a large corpus by an order of magnitude after integrating hardware acceleration using Connectal. Performance improvements on the Zynq-based platforms was even greater due to the relative processing power of the ARM CPU and scaled with the number of bus master interfaced used for DMA. In the Connectal framework, developing these applications took very little time.

Performance of Portals

The current implementation of HW/SW textbf{portal} transfers 32 bits per FPGA clock cycle. Our example designs run at 100MHz to 250MHz, depending on the complexity of the design and the speed grade of the FPGA used. Due to their intended use, the important performance metric of Portals is latency. These values are given in Figure~ref{Fig:PortalLatency}.

```
begin{figure} centering begin{tabular}{|l|}
\hline & rt{KC705} & rt{VC707} & rt{ZYBO} & rt{Zedboard} & rt{ZC702} &
rt{ZC706} & rt{Parallel} & rt{Mini-ITX} \
\hline HW  $\rightarrow$  SW & 3 & 3 & X & 0.80 & 0.80 & 0.65 & X & 0.65 \ hline SW
 $\rightarrow$  HW & 5 & 5 & X & 1.50 & 1.50 & 1.10 & X & 1.10 \ hline
end{tabular} caption{Latency ( $\mu$ s) of communication through portals on supported
platformslabel{Fig:PortalLatency}}
end{figure}
```

The Xilinx KC705 and VC707 boards connect to x86 CPUs and system memory via PCIe gen1. The default FPGA clock for those boards is 125MHz. The other platforms use AXI to connect the programmable logic to the quad-core ARM Cortex A9 and system memory. The ZYBO, Zedboard and ZC702 use a slower speed grade part on which our designs run at 100MHz. The ZC706 and Mini-ITX use a faster part on which many of our designs run at 200MHz. The lower latency measured on the ZC706 reflects the higher clock speed of the latency performance test.

Performance of Reads/Writes of System Memory

For high bandwidth transfers, we assume the developer will have the application hardware read or write system memory directly. Direct access to memory enables transfers with longer bursts, reducing memory bus protocol overhead.

The framework supports transfer widths of 32 to 128 bits per cycle, depending on the interconnect used.

Our goal in the design of the library components used to read and write system memory is to ensure that a developer's application can use all bandwidth available to the FPGA when accessing system memory. DMA Bandwidth on supported platforms is listed in [Figure 16.10](#).

```
begin{figure} centering begin{tabular}{cccccccc}
\hline & rt{KC705} & & rt{VC707} & & rt{ZYBO} & & rt{Zedboard} & & rt{ZC702} & &
rt{ZC706} & & rt{Parallel} & & rt{Mini-ITX} & \
\\
\hline Read & 1.4 & & 1.4 & & X & & 0.8 & & 0.8 & & 1.6 & & X & & 1.6 \\
\hline Write & 1.4 & & 1.4 & & X & & 0.8 & & 0.8 & & 1.6 & & X & & 1.6 \\
\hline
\end{tabular} caption{Maximum bandwidth (GB/s) between FPGA and host memory using
Connectal RTL libraries on supported platforms\label{Fig:DmaBandwidth}}
end{figure}
```

On PCIe systems, Connectal currently supports 8 lane PCIe gen1. We've measured 1.4 gigabytes per second for both reads and writes. Maximum throughput of 8 lane PCIe gen1 is 1.8GB/s, taking into account 1 header transaction per 8 data transactions, where 8 is the maximum number of data transactions per request supported by our server's chipset. The current version of the test needs some more tuning in order to reach the full bandwidth available. In addition, we are in the process of updating to 8 lane PCIe gen2 using newer Xilinx cores.

Zynq systems have four *high performance* ports for accessing system memory. Connectal enables an accelerator to use all four. In our experiments, we have been able to achieve 3.6x higher bandwidth using 4 ports than using 1 port.

Related Work

A number of research projects, such as Lime~cite{REL:Lime}, BCL~cite{King:2012:AGH:2150976.2151011}, HThreads~cite{DBLP:conf/fpl/PeckAASBA06}, and CatapultC~cite{CatC:www} (to name just a few) bridge the software/hardware development gap by providing a single language for developing both the software and hardware components of the design. In addition, Altera and Xilinx have both implemented OpenCL~cite{Opencl} on FPGAs~cite{AlteraOpencl,XilinxOpencl} in an attempt to attract GPU programmers.

The computation model of software differs significantly from that of hardware, and so far none of the unified language approaches deliver the same performance as languages designed specifically for hardware or software. Connectal is intended to be used for the design of performance-critical systems. In this context we think that designers prefer a mix of languages specifically designed for their respective implementation contexts.

Infrastructures such as LEAP~cite{DBLP:conf/fpl/FlemingYAE14}, Rainbow~cite{DBLP:journals/ijrc/JozwikHETT13}, and OmpSs~cite{DBLP:conf/fpga/FilguerasGJAMLNV14} (to name just a few) use resource abstraction to enable FPGA development. We found that in their intended context, these tools were easy to use but that performance tuning in applications not foreseen by the infrastructure developers was problematic.

Some projects such as TMD-MPI~cite{DBLP:journals/trets/SaldanaPMNWCWSP10}, VFORCE/VSIPL++~cite{DBLP:journals/jpdc/MooreLK12}, and GASNet/GAScore~cite{DBLP:conf/fpga/WillenbergC13} target only the hardware software interface. These tools provide message passing capabilities, but rely on purely operational semantics to describe the HW/SW interface. Apart from the implementation details, Connectal distinguishes itself by using an IDL to enforce denotational interface semantics.

UIO~cite{UIO:Howto} is a user-space device driver framework for Linux. It is very similar to the Connectal's portal device driver, but it does not provide a solution to multiple device nodes per hardware device. The portal driver provides this so that different interfaces of a design may be accessed independently, providing process boundary protection, thread safety, and the ability for user processes and the kernel both to access the hardware device.

Conclusion

Connectal bridges the gap between software and hardware development, enabling developers to create integrated solutions rapidly. With Connectal, we take a pragmatic approach to software and hardware development in which we try to avoid any dependence on proposed solutions to open research problems.

Use of Connectal's interface compiler ensures that software and hardware remain consistent and make it easy to update the hardware/software boundary as needed in a variety of execution contexts. The generated portals permit concurrent and low-latency access to the accelerator and enable different processes or the kernel to have safe isolated access through dedicated interfaces. Support for sharing memory between software and hardware makes it easy to achieve high transfer speeds between the two environments.

Connectal supports Linux and Android operating systems running on x86 and ARM CPUs. It currently supports Xilinx FPGAs and runs on the full range of Xilinx Series 7 devices. Our fully-scripted development process enables the use of continuous integration of software and hardware development. Integrating software development early makes it easier to ensure that the complete solution actually meets design targets and customer requirements.

Portal Interface Structure

Connectal connects software and hardware via portals, where each portal is an interface that allows one side to invoke methods on the other side.

We generally call a portal from software to hardware to be a “request” and from hardware to software to be an “indication” interface:

```
Sequence Diagram to be drawn
{
    SW; HW
    SW -> HW [label = "request"];
    SW <- HW [label = "indication"];
}
```

A portal is conceptually a FIFO, where the arguments to a method are packaged as a message. **CONNECTAL** generates a “proxy” that marshalls the arguments to the method into a message and a “wrapper” that unpacks the arguments and invokes the method.

Currently, connectal Includes a library that implements portals from software to hardware via memory mapped hardware FIFOs.

Portal Device Drivers

Connectal uses a platform-specific driver to enable user-space applications to memory-map each portal used by the application and to enable the application to wait for interrupts from the hardware.

indexterm:pcieportal indexterm:zynqportal

- pcieportal.ko
- zynqportal.ko

Connectal also uses a generic driver to enable the applications to allocate DRAM that will be shared with the hardware and to send the memory mapping of that memory to the hardware.

- portalmem.ko

Portal Memory Map

Connectal is designed to support multiple tiles, each of which can hold an independent design. Currently, the number of tiles is one.

Connectal currently supports up to 16 portals connected between software and hardware, for a total of 64KB of address space.

Base address	Function
0x0000	Portal 0
0x1000	Portal 1
0x2000	Portal 2
0x3000	Portal 3
0x4000	Portal 4
0x5000	Portal 5
0x6000	Portal 6
0x7000	Portal 7
0x8000	Portal 8
0x9000	Portal 9
0xa000	Portal 10
0xb000	Portal 11
0xc000	Portal 12
0xd000	Portal 13
0xe000	Portal 14
0xf000	Portal 15

Each portal uses 16KB of address space, consisting of a control register region and then per-method FIFOs, each of which takes 32 bytes of address space.

Base address	Function
0x000	Portal control regs
0x020	Method 0 FIFO
0x040	Method 1 FIFO
...	...

For request portals, the FIFOs are from software to hardware, and for indication portals the FIFOs are from hardware to software.

Portal FIFOs

Base address	Function
0x00	FIFO data (write request data, read indication data)
0x04	Request FIFO not full / Indication FIFO not empty

Portal Control Registers

Base address	Function	Description
0x00	Interrupt status register	1 if this portal has any messages ready, 0 otherwise
0x04	Interrupt enable register	Write 1 to enable interrupts, 0 to disable
0x08	Number of tiles	
0x0C	Ready Channel number + 1	Reads as zero if no indication channel ready
0x10	Interface Id	
0x14	Number of portals	
0x18	Cycle count LSW	Snapshots MSW when read
0x1C	Cycle count MSW	MSW of cycle count when LSW was read

What is Connectal?

Connectal provides a hardware-software interface for applications split between user mode code and custom hardware in an FPGA or ASIC.

Connectal can automatically build the software and hardware glue for a message based interface and also provides for configuring and using shared memory between applications and hardware. Communications between hardware and software are provided by a bidirectional flow of events and regions of memory shared between hardware and software. Events from software to hardware are called requests and events from hardware to software are called indications, but in fact they are symmetric.

Lexicon

connectal:: The name of the project, whose goal is to ease the task of building applications composed of hardware and software components. Programmers use bsv as an IDL to specify the interface between the hardware and software components. A combination of generated code and libraries coordinate the data-flow between the program modules. Because the HW and SW stacks are customized for each application, the overheads associated with communicating across the HW/SW boundary are low.

HW/SW interface :: portal

bsv:: Bluespec System Verilog. bsv is a language for describing hardware that is might higher level than verilog. See {bsvdocumentation}[BSV Documentation] and {bluespecdotcom}[Bluespec, Inc].

bluespec:: Shorthand for Bluespec System Verilog (bsv)

indexterm:portal portal:: a logical request/indication pair is referred to as a portal. current tools require their specification in the IDL to be syntactically identifiable (i.e. fooRequest/fooIndication). An application can make use of multiple portals, which may be specified independently.

request interface:: These methods are implemented by the application hardware to be invoked by application software. A bsv interface consisting of ‘Action’ methods. Because of the ‘Action’ type, data flow across this interface is unidirectional (SW -> HW).

indication interface:: The dual of a request interface, indication interfaces are ‘Action’ methods implemented by application software to be invoked by application hardware. As with request interfaces, the data flow across this interface is unidirectional, but in the opposite direction.

pcieportal/zynqportal:: these two loadable kernel modules implement the minimal set of driver functionality. Specifically, they expose portal HW registers to SW through mmap, and set up interrupts to notify SW that an indication method has been invoked by HW.

portalalloc:: This loadable kernel module exposes a subset of dma-buf functionality to user-space software (though a set of ioctl commands) to allocate and manage memory regions which can be shared between SW and HW processes.

Maintaining coherence of the allocated buffers between processes is not automatic: ioctl commands for flush/invalidate are provided to be invoked explicitly by the users if necessary.

connectalgen:: The name of the interface compiler which takes as input the bsv interface specification along with a description of a target platform and generates logic in both HW and SW to support this interface across the communication fabric.

Example setups:

A zedboard (<http://www.zedboard.org/>), with Android running on the embedded ARM processors (the Processing System 7), an application running as a user process, and custom hardware configured into the Programmable Logic FPGA.

An x86 server, with Linux running on the host processor, an application running partly as a user process on the host and partly as hardware configured into an FPGA connected by PCI express (such as the Xilinx VC707 (<http://www.xilinx.com/products/boards-and-kits/EK-V7-VC707-G.htm>)).

Background

When running part or all of an application in an FPGA, it is usually necessary to communicate between code running in user mode on the host and the hardware. Typically this has been accomplished by custom device drivers in the OS, or by shared memory mapped between the software and the hardware, or both. Shared memory has been particularly troublesome under Linux or Android, because devices frequently require contiguous memory, and the mechanisms for guaranteeing successful memory allocation often require reserving the maximum amount of memory at boot time.

Portal tries to provide convenient solutions to these problems in a portable way.

It is desirable to have

- low latency for small messages
- high bandwidth for large messages
- notification of arriving messages
- asynchronous replies to messages
- support for hardware simulation by a separate user mode process
- support for shared memory (DMA) between hardware and software

Overview

Portal is implemented as a loadable kernel module device driver for Linux/Android and a set of tools to automatically construct the hardware and software glue necessary for communications.

Short messages are handled by programmed I/O. The message interface from software to hardware (so called “requests”) is defined as a bsv interface containing a number of Action methods, each with a name and typed arguments. The interface generator creates all the software and hardware glue so that software invocations of the interface stubs flow through to, and are turned into bsv invocations of the matching hardware. The machinery does not have flow control. Software is responsible for not overrunning the hardware. There is a debug mechanism which will return the request type of a failed method, but it does not tell which invocation failed. Hardware to software interfaces (so called “indications”) are likewise defined by bsv interfaces containing Action methods. Hardware invocations of these methods flow through to and cause software calls to corresponding user-supplied functions. In the current implementation there is flow control, in that the hardware will stall until there is room for a hardware to software message. There is also a mechanism for software to report a failure, and there is machinery for these failures to be returned to the hardware.


```
{ // edge label SW -> HW [label = "request"]; SW <- HW [label = "indication"];
```

Portals do not have to be structured as request/response. Hardware can send messages to software without a prior request from software.

```
{ // edge label SW <- HW [label = "indication"];
```

Incoming messages can cause host interrupts, which wake up the device driver, which can wake up the user mode application by using the `select(2)` or `poll(2)` interfaces.

Most of the time, communications between hardware and software will proceed without requiring use of the OS. User code will read and write directly to memory mapped I/O space. Library code will poll for incoming messages, and [true? eventually time out and call `poll(2)`. Only when `poll(2)` or `select(2)` are called will the device driver enable hardware interrupts. Thus interrupts are only used to wake up software after a quiet period.

The designer specifies a set of hardware functions that can be called from software, and a set of actions that the hardware can take which result in messages to software. Portal tools take this specification and build software glue modules to translate software function calls into I/O writes to hardware registers, and to report hardware events to software.

For larger memory and OS bypass (OS bypass means letting the user mode application talk directly to the hardware without using the OS except for setup), portal implements shared memory. Portal memory objects are allocated by the user mode program, and appear as Linux file descriptors. The user can `mmap(2)` the file to obtain user mode access to the shared memory region. Portal does not assure that the memory is physically contiguous, but does pin it to prevent the OS from reusing the memory. An FPGA DMA controller module is provided that gives the illusion of contiguous memory to application hardware, while under the covers using a translation table of scattered addresses.

The physical addresses are provided to the user code in order to initialize the dma controller, and address “handles” are provided for the application hardware to use.

The DMA controller provides Bluespec objects that support streaming access with automatic page crossings, or random access.

An Example

An application developer will typically write the hardware part of the application in Bluespec and the software part of the application in C or C++. In a short example, there will be a bsv source file for the hardware and a cpp source file for the application.

The application developer is free to specify whatever hardware-software interface makes sense.

Refer to <https://github.com/cambridgehackers/connectal>

In the examples directory, see [simple](../examples/simple/). The file [Simple.bsv](../examples/simple/Simple.bsv) defines the hardware, and `testsimple.cpp` supplies the software part. In this case, the software part is a test framework for the hardware.

Simple.bsv declares a few *struct* and *enum* types:

```
typedef struct{
    Bit#(32) a;
    Bit#(32) b;
} S1 deriving (Bits);

typedef struct{
    Bit#(32) a;
    Bit#(16) b;
    Bit#(7) c;
} S2 deriving (Bits);
```

```
typedef enum {
    ElChoice1,
    ElChoice2,
    ElChoice3
} El deriving (Bits,Eq);

typedef struct{
    Bit#(32) a;
    El e1;
} S3 deriving (Bits);
```

Simple.bsv defines the actions (called Requests) that software can use to cause the hardware to act, and defines the notifications (called Indications) that the hardware can use to signal the software.

```
interface SimpleIndication;
    method Action heard1(Bit#(32) v);
    method Action heard2(Bit#(16) a, Bit#(16) b);
    method Action heard3(S1 v);
    method Action heard4(S2 v);
    method Action heard5(Bit#(32) a, Bit#(64) b, Bit#(32) c);
    method Action heard6(Bit#(32) a, Bit#(40) b, Bit#(32) c);
    method Action heard7(Bit#(32) a, El e1);
endinterface

interface SimpleRequest;
    method Action say1(Bit#(32) v);
    method Action say2(Bit#(16) a, Bit#(16) b);
    method Action say3(S1 v);
    method Action say4(S2 v);
    method Action say5(Bit#(32) a, Bit#(64) b, Bit#(32) c);
    method Action say6(Bit#(32) a, Bit#(40) b, Bit#(32) c);
    method Action say7(S3 v);
endinterface
```

Software can start the hardware working via say, say2, ... Hardware signals back to software with heard and heard2 and so fort. In the case of this example, say and say2 merely echo their arguments back to software.

The definitions in the bsv file are used by the connectal infrastructure (a python program) to automatically create corresponding c++ interfaces.:

```
../../connectalgen -Bbluesim -p bluesim -x mkBsimTop \
-s2h SimpleRequest \
-h2s SimpleIndication \
-s testsimple.cpp \
-t ../../bsv/BsimTop.bsv Simple.bsv Top.bsv
```

The tools have to be told which interface records should be used for Software to Hardware messages and which should be used for Hardware to Software messages. These interfaces are given on the command line for genxpprojfrombsv

connectalgen constructs all the hardware and software modules needed to wire up portals. This is sort of like an RPC compiler for the hardware-software interface. However, unlike an RPC each method is asynchronous.

The user must also create a toplevel bsv module Top.bsv, which instantiates the user portals, the standard hardware environment, and any additional hardware modules.

Rather than constructing the *makefilegen* command line from scratch, the examples in connectal use include *Makefile.connectal* and define some *make* variables.

Here is the Makefile for the *simple* example:

```
CONNECTALDIR?=/../..
INTERFACES = SimpleRequest SimpleIndication
BSVFILES = Simple.bsv Top.bsv
CPPFILES=testsimple.cpp

include $(CONNECTALDIR)/Makefile.connectal
```

Designs outside the connectal directory using *connectal* may also include *Makefile.connectal*:

```
CONNECTALDIR=/scratch/connectal
INTERFACES = ...
BSVFILES = ...
CPPFILES = ...
include $(CONNECTALDIR)/Makefile.connectal
```

simple/Top.bsv

Each CONNECTAL design implements [Top.bsv](../examples/simple/Top.bsv) with some standard components.

It defines the *IfcNames* enum, for use in identifying the portals between software and hardware:

```
typedef enum {SimpleIndication, SimpleRequest} IfcNames deriving (Eq,Bits);
```

It defines *mkConnectalTop*, which instantiates the wrappers, proxies, and the design itself:

```
module mkConnectalTop (StdConnectalTop# (addrWidth)) ;
```

:bsv:module:StdConnectalTop is parameterized by *addrWidth* because Zynq and x86 have different width addressing. *StdConnectalTop* is a typedef:

```
typedef ConnectalTop# (addrWidth, 64, Empty) StdConnectalTop# (numeric type addrWidth);
```

The “64” specifies the data width and *Empty* specifies the empty interface is exposed as pins from the design. In designs using HDMI, for example, *Empty* is replaced by *HDMI*. On some platforms, the design may be able to use different data widths, such as 128 bits on x86/PCIe.

Next, *mkConnectalTop* instantiates user portals:

```
// instantiate user portals
SimpleIndicationProxy simpleIndicationProxy <- mkSimpleIndicationProxy (SimpleIndication);
```

Instantiate the design:

```
SimpleRequest simpleRequest <- mkSimpleRequest (simpleIndicationProxy.ifc);
```

Instantiate the wrapper for the design:

```
SimpleRequestWrapper simpleRequestWrapper <- mkSimpleRequestWrapper (SimpleRequest, simpleRequest);
```

Collect the portals into a vector:

```
Vector#(2, StdPortal) portals;
portals[0] = simpleRequestWrapper.portalIfc;
portals[1] = simpleIndicationProxy.portalIfc;
```

Create an interrupt multiplexer from the vector of portals:

```
let interrupt_mux <- mkInterruptMux (portals);
```

Create the system directory, which is used by software to locate each portal via the *IfcNames* enum:

```
// instantiate system directory
StdDirectory dir <- mkStdDirectory(portals);
let ctrl_mux <- mkAxisSlaveMux(dir,portals);
```

The following generic interfaces are used by the platform specific top BSV module:

```
interface interrupt = interrupt_mux;
interface ctrl = ctrl_mux;
interface m_axi = null_axi_master;
interface leds = echoRequestInternal.leds;

endmodule : mkConnectalTop
```

simple/testsimple.cpp

CONNECTAL generates header files declaring wrappers for hardware-to-software interfaces and proxies for software-to-hardware interfaces. These will be in the “jni/” subdirectory of the project directory.

```
#include "SimpleIndication.h"
#include "SimpleRequest.h"
```

It also declares software equivalents for structs and enums declared in the processed BSV files:

```
#include "GeneratedTypes.h"
```

CONNECTAL generates abstract virtual base classes for each Indication interface.

```
class SimpleIndicationWrapper : public Portal {
public:
    ...
    SimpleIndicationWrapper(int id, PortalPoller *poller = 0);
    virtual void heard1 ( const uint32_t v )= 0;
    ...
};
```

Implement subclasses of the wrapper in order to define the callbacks:

```
class SimpleIndication : public SimpleIndicationWrapper
{
public:
    ...
    virtual void heard1(uint32_t a) {
        fprintf(stderr, "heard1(%d)\n", a);
        assert(a == vla);
        incr_cnt();
    }
    ...
};
```

To connect these classes to the hardware, instantiate them using the *IfcNames* enum identifiers. CONNECTAL prepends the name of the type because C++ does not support overloading of enum tags.

```
SimpleIndication *indication = new SimpleIndication(IfcNames_SimpleIndication);
SimpleRequestProxy *device = new SimpleRequestProxy(IfcNames_SimpleRequest);
```

Create a thread for handling notifications from hardware:

```
pthread_t tid;
if(pthread_create(&tid, NULL, portalExec, NULL)){
    exit(1);
}
```

Now the software invokes hardware methods via the proxy:

```
device->say1(v1a);

device->say2(v2a,v2b);
```

Simple Example Design Structure

The *simple* example consists of the following files:

```
Simple.bsv
Makefile
Top.bsv
testsimple.cpp
```

After running *make BOARD=zedboard verilog* in the *simple* directory, the *zedboard* project directory is created, populated by the generated files.

A top level *Makefile* is created:

```
zedboard/Makefile
```

makefilegen generates wrappers for software-to-hardware interfaces and proxies for hardware-to-software interfaces:

```
zedboard/sources/mkzynqtop/SimpleIndicationProxy.bsv
zedboard/sources/mkzynqtop/SimpleRequestWrapper.bsv
```

Connectal supports Android on Zynq platforms, so connectalgen generates *jni/Android.mk* for *ndk-build*:

```
zedboard/jni/Android.mk
zedboard/jni/Application.mk
```

Connectal generates *jni/Makefile* to compile the software for PCIe platforms (vc707 and kc705):

```
zedboard/jni/Makefile
```

CONNECTAL generates software proxies for software-to-hardware interfaces and software wrappers for hardware-to-software interfaces:

```
zedboard/jni/SimpleIndication.h
zedboard/jni/SimpleIndication.cpp
zedboard/jni/SimpleRequest.cpp
zedboard/jni/SimpleRequest.h
```

CONNECTAL also generates *GeneratedTypes.h* for struct and enum types in the processed BSV source files:

```
zedboard/jni/GeneratedTypes.h
```

CONNECTAL copies in standard and specified constraints files:

```
zedboard/constraints/design_1_processing_system7_1_0.xdc
zedboard/constraints/zedboard.xdc
```

CONNECTAL generates several TCL files to run *vivado*.

The *board.tcl* file specifies *partname*, *boardname*, and *connectaldir* for the other TCL scripts.:

```
zedboard/board.tcl
```

To generate an FPGA bit file, run *make bits*. This runs vivado with the *mkzynqtop-impl.tcl* script.:

```
zedboard/mkzynqtop-impl.tcl
```

make verilog

Compiling to verilog results in the following verilog files:

```
zedboard/verilog/top/mkSimpleIndicationProxySynth.v
zedboard/verilog/top/mkZynqTop.v
```

Verilog library files referenced in the design are copied for use in synthesis.:

```
zedboard/verilog/top/FIFO1.v
...
```

make bits

Running *make bits* in the zedboard directory results in timing reports:

```
zedboard/bin/mkzynqtop_post_place_timing_summary.rpt
zedboard/bin/mkzynqtop_post_route_timing_summary.rpt
zedboard/bin/mkzynqtop_post_route_timing.rpt
```

and some design checkpoints:

```
zedboard/hw/mkzynqtop_post_synth.dcp
zedboard/hw/mkzynqtop_post_place.dcp
zedboard/hw/mkzynqtop_post_route.dcp
```

and the FPGA configuration file in .bit and .bin formats:

```
zedboard/hw/mkZynqTop.bit
zedboard/hw/mkZynqTop.bin
```

make android_exe

CONNECTAL supports Android 4.0 on Zynq platforms. It generates *jni/Android.mk* which is used by *ndk-build* to create a native Android executable.:

```
make android_exe
```

This produces the ARM elf executable:

```
libs/armeabi/android_exe
```

make run

For Zynq platforms:

```
make run
```

will copy the Android executable and FPGA configuration file to the target device, program the FPGA, and run the executable. See [run.android](../scripts/run.android) for details.

It uses *checkip* to determine the IP address of the device via a USB console connection to the device (it is built/installed on the host machine from the git repo [cambridgehackers/consolable](#)). If the target is not connected to the build machine via USB, specify the IP address of the target manually:

```
make RUNPARAM=ipaddr run
```

For PCIe platforms, *make run* programs the FPGA via USB and runs the software locally.

For bluesim, *make run* invokes bluesim on the design and runs the software locally.

Shared Memory

Shared Memory Hardware

In order to use shared memory, the hardware design instantiates a DMA module in Top.bsv:

```
AxiDmaServer#(addrWidth,64) dma <- mkAxiDmaServer(dmaIndicationProxy.ifc, readClients, writeClients),
```

The AxiDmaServer multiplexes read and write requests from the clients, translates DMA addresses to physical addresses, initiates bus transactions to memory, and delivers responses to the clients.

DMA requests are specified with respect to “portal” memory allocated by software and identified by a *pointer*.

Requests and responses are tagged in order to enable pipelining:

```
typedef struct {
    SGLId pointer;
    Bit#(MemOffsetSize) offset;
    Bit#(8) burstLen;
    Bit#(6) tag;
} MemRequest deriving (Bits);

typedef struct {
    Bit#(dsz) data;
    Bit#(6) tag;
} MemData#(numeric type dsz) deriving (Bits);
```

Read clients implement the *MemReadClient* interface. On response to the read, *burstLen* *MemData* items will be put to the *readData* interface. The design must be ready to consume the data when it is delivered from the memory bus or the system may hang:

```
interface MemReadClient#(numeric type dsz);
    interface GetF#(MemRequest)    readReq;
    interface PutF#(MemData#(dsz)) readData;
endinterface
```

Write clients implement *MemWriteClient*. To complete the transaction, *burstLen* data items will be consumed from the *writeData* interface. Upon completion of the request, the specified tag will be put to the *writeDone* interface. The data must be available when the write request is issued to the memory bus or the system may hang:

```
interface MemWriteClient#(numeric type dsz);
    interface GetF#(MemRequest)    writeReq;
    interface GetF#(MemData#(dsz)) writeData;
    interface PutF#(Bit#(6))       writeDone;
endinterface
```

A design may implement *MemReadClient* and *MemWriteClient* interfaces directly, or it may instantiate *DmaReadBuffer* or *DmaWriteBuffer*.

The *AxiDmaServer* is configured with physical address translations for each region of memory identified by a *pointer*. A design using DMA must export the *DmaConfig* and *DmaIndication* interfaces of the DMA server.

Here are the DMA components of [memread_nobuff/Top.bsv](../examples/memread_nobuff/Top.bsv):

Instantiate the design and its interface wrappers and proxies:

```
MemreadIndicationProxy memreadIndicationProxy <- mkMemreadIndicationProxy (MemreadIndication);
Memread memread <- mkMemread (memreadIndicationProxy.ifc);
MemreadRequestWrapper memreadRequestWrapper <- mkMemreadRequestWrapper (MemreadRequest, memread.request);
```

Collect the read and write clients:

```
Vector#(1, MemReadClient#(64)) readClients = cons (memread.dmaClient, nil);
Vector#(0, MemReadClient#(64)) writeClients = nil;
```

Instantiate the DMA server and its wrapper and proxy:

```
DmaIndicationProxy dmaIndicationProxy <- mkDmaIndicationProxy (DmaIndication);
AxiDmaServer#(addrWidth, 64) dma <- mkAxiDmaServer (dmaIndicationProxy.ifc, readClients, writeClients);
DmaConfigWrapper dmaConfigWrapper <- mkDmaConfigWrapper (DmaConfig, dma.request);
```

Include *DmaConfig* and *DmaIndication* in the portals of the design:

```
Vector#(4, StdPortal) portals;
portals[0] = memreadRequestWrapper.portalIfc;
portals[1] = memreadIndicationProxy.portalIfc;
portals[2] = dmaConfigWrapper.portalIfc;
portals[3] = dmaIndicationProxy.portalIfc;
```

The code generation tools will then produce the software glue necessary for the shared memory support libraries to initialize the DMA “library module” included in the hardware.

Shared Memory Software

The software side instantiates the *DmaConfig* proxy and the *DmaIndication* wrapper:

```
dma = new DmaConfigProxy (IfcNames_DmaConfig);
dmaIndication = new DmaIndication (dma, IfcNames_DmaIndication);
```

Call *dma->alloc()* to allocate DMA memory. Each chunk of portal memory is identified by a file descriptor. Portal memory may be shared with other processes. Portal memory is reference counted according to the number of file descriptors associated with it:

```
PortalAlloc *srcAlloc;
dma->alloc (alloc_sz, &srcAlloc);
```

Memory map it to make it accessible to software:

```
srcBuffer = (unsigned int *)mmap (0, alloc_sz, PROT_READ|PROT_WRITE|PROT_EXEC, MAP_SHARED, srcAlloc->fd);
```

Connectal is currently using non-snooped interfaces, so the cache must be flushed and invalidated before hardware accesses portal memory:

```
dma->dCacheFlushInval (srcAlloc, srcBuffer);
```

Call *dma->reference()* to get a pointer that may be passed to hardware:


```
unsigned int ref_srcAlloc = dma->reference(srcAlloc);
```

This also transfers the DMA-to-physical address translation information to the hardware via the *DmaConfig* interface:

```
device->startRead(ref_srcAlloc, numWords, burstLen, iterCnt);
```

Notes

stewart notes

Currently there are no valid bits and no protections against bursts crossing page boundaries

There needs to be a way to synchronize Request actions and DMA reads, and to synchronize DMA writes with Indications, so that the writes complete to the coherence point before the indication is delivered to software. One could imagine an absurdly buffered memory interface and a rather direct path for I/O reads that could get out of order.

Interface Declarations

Flow Control

Host Interface

Connectal Developer's Guide

Connectal Rationale

Are you happy with the connection between software development and hardware development on your projects? Do you wish you had more flexible tools for developing test benches, drivers, and applications using your hardware? Do you wish you didn't need a kernel driver?

Do you wish that you could use a common flow to drive your hardware development, across simulation, timing verification, FPGA implementation, and even the ASIC back end?

Do you wish you could use a better programming language to implement the hardware? Do you miss declared interfaces and structural types?

Do you wish you had better tools for debugging the hardware? For reflecting conditions encountered on the hardware back into the simulation environment where you can drill down into the problem?

These are the kinds of problems we encountered which inspired us to develop the Connectal framework. With Connectal, we are trying to bring the productivity of software engineering practices to hardware development without compromising on the performance of the hardware.

Connectal provides:

- Seamless development environment with test benches and apps written in C/C++
- Streamlined application structure
- Faster builds and greater visibility simulating hardware in bluesim/modelsim/xsim
- Support for capture traces on interfaces and replaying them in test benches

Connectal Project Structure

The set of files composing the input to the Connectal toolchain is referred to as a project. A collection of out-of-tree example projects is available at <https://github.com/connectal-examples>. To illustrate the structure of a project, this chapter uses the example <https://github.com/connectal-examples/leds>, which can be executed on the Bluesim or Zynq target platforms.

Project Makefile

The top-level Makefile in the project (in our example, <https://github.com/connectal-examples/leds/blob/master/Makefile>) defines parameters building and executing the project. In its simplest form, it specifies the hardware and software source files, which Bluespec interfaces to use as interfaces (portals) between them, and the libraries to use for the hardware and software compilation:

```
INTERFACES = LedControllerRequest
BSVFILES = LedController.bsv Top.bsv
CPPFILES= testleds.cpp
include \$(CONNECTALDIR)/Makefile.connectal
```

BSVFILES is a list of bsv files containing interface definitions used to generate portals and module definitions used to generate HW components. Connectal bsv libraries can be used without being listed explicitly.

CPPFILES is a list of C/C++ files containing software components and main. The Connectal C/C++ libraries can be used without being listed explicitly.

INTERFACES is a list of names of BSV interfaces which may be used to communicate between the HW and SW components. In addition to user-defined interfaces, there are a wide variety of interfaces defined in Connectal libraries which may be included in this list.

NUMBER_OF_MASTERS is used to designate the number of host bus masters the hardware components will instantiate. For PCIe-based platforms, this value can be set to 1, while on Zynq-based platforms values from 1 to 4 are valid.

CONNECTALDIR must be set so that the top-level Connectal makefile can be included, defining the makefile build targets for the project. This brings in the default definitions of all project build parameters as well as the Connectal hardware and software libraries. When running the toolchain on AWS, this variable is set automatically in the build environment. (See *Compiling and Running Connectal Project*)

Project Source

Interface Definitions

```
label{interface_definitions}
```

When generating portals, the Connectal interface compiler searches the Connectal bsv libraries and the files listed in BSVFILES for definitions of the interfaces listed in INTERFACES. If the definition of a listed interfaces is not found, an error is reported and the compilation aborts. The interfaces in this list must be composed exclusively of Action methods. Supported method argument types are Bit\#(n), Bool, Int\#(32), UInt\#(32), Float, Vector\#(t), enum, and struct.

Software

The software in a Connectal project consists of at least one C++ file which instantiates the generated portal wrappers and proxies and implements main(). The following source defines the SW component of the example, which simply toggles LEDs on the Zedboard ([url{https://github.com/connectal-examples/leds/blob/master/testleds.cpp}](https://github.com/connectal-examples/leds/blob/master/testleds.cpp)):

```

#include <unistd.h>
#include "LedControllerRequest.h"
#include "GeneratedTypes.h"
int main(int argc, const char **argv)
{
    LedControllerRequestProxy *device =
        new LedControllerRequestProxy(IfcNames_LedControllerRequest);
    for (int i = 0; i < 20; i++) {
        device->setLeds(10, 10000);
        sleep(1);
        device->setLeds(5, 10000);
        sleep(1);
    }
}

```

The makefile listed `LedControllerRequest` as the only communication interface. The generated proxies and wrappers for this interface are in `LedControllerRequest.h` which is included, along with C++ implementations of all additional interface types in `GeneratedTypes.h`. Line 9 instantiates the proxy through which the software invokes the hardware methods (See also *Flow Control*)

To support projects that will execute software inside the linux kernel (for example, to work in conjunction with filesystems or the network stack), connectal project software can also be written in C.

Hardware

Connectal projects typically have at least one BSV file containing interface declarations and module definitions. The implementation of the interfaces and all supporting infrastructure is standard BSV. Interfaces being used as portals are subject to the type restrictions described earlier (See also *Interface Declarations*)

Top.bsv

In `Top.bsv` (<https://github.com/connectal-examples/leds/blob/master/Top.bsv>), the developer instantiates all hardware modules explicitly. Interfaces which can be invoked through portals need to be connected to the generated wrappers and proxies. To connect to the host processor bus, a parameterized standard interface is used, making it easy to synthesize the application for different CPUs or for simulation:

```

// Connectal Libraries
import CtrlMux::*;
import Portal::*;
import Leds::*;
import MemTypes::*;
import MemPortal::*;
import HostInterface::*;
import LedControllerRequest::*;
import LedController::*;

typedef enum {LedControllerRequestPortal} IfcNames deriving (Eq,Bits);

module mkConnectalTop(StdConnectalTop#(PhysAddrWidth));
    LedController ledController <- mkLedControllerRequest();
    LedControllerRequestWrapper ledControllerRequestWrapper <-
        mkLedControllerRequestWrapper(LedControllerRequestPortal,
            ledController.request);

    Vector#(1,StdPortal) portals;
    portals[0] = ledControllerRequestWrapper.portalIfc;

```

```
let ctrl_mux <- mkSlaveMux(portals);

interface interrupt = getInterruptVector(portals);
interface slave = ctrl_mux;
interface masters = nil;
interface leds = ledController.leds;
endmodule
```

Like the SW components, the HW begins by importing the generated wrappers and proxies corresponding to the interfaces listed in the project Makefile. The user-defined implementation of the `LedControllerRequest` interface is instantiated on line 14, and wrapped on line 15. This wrapped interface is connected to the bus using the library module `mkSlaveMux` on line 21 so it can be invoked from the software. At the end of the module definition, the top-level interface elements must be connected. A board-specific top-level module will include this file, instantiate `mkConnectalTop` and connect the interfaces to the actual peripherals. The module `mkConnectalTop` must be defined in a file named `Top.bsv` in the user project.

The Bluespec compiler generates a Verilog module from the top level BSV module, in which the methods of exposed interfaces are implemented as Verilog ports. Those ports are bound to physical pins on the FPGA using a physical constraints file. If CPU specific interface signals are needed by the design (for example, extra clocks that are generated by the PCIe core), then an optional CPU-specific `HostInterface` parameter to `mkConnectalTop` can also be used. If the design uses external pins on the FPGA, those connections are also made here by exporting a ‘Pins’ interface ([hyperref\[host_interface\]](#){[Section~ref{host_interface}](#)}) and providing bindings in the constraints file.

Compiling and Running Connectal Project

Compiling on ConnectalBuild

The Connectal toolchain can be run on ConnectalBuild using the following Buildbot web interface: <http://connectalbuild.qrclab.com/projects>.

The screenshot shows a web browser window with the address `connectalbuild.qrclab.com/projects`. The page has a navigation bar with links: Home, Waterfall, Grid, T-Grid, Console, Builders, Recent Builds, Buildslaves, Changesources, JSON API, and About. A login button is also present. The main heading reads "Welcome to the Connectal Buildbot for the Connectal project!". Below this is a form with the following fields: "Project Name:" (text input), "Project Repo:" (text input with the value `git://github.com/<username>/<repo>.git`), "Path:" (text input), "Branch:" (text input), "Revision:" (text input), "Bluesim:" (checkbox), "Zedboard:" (checkbox), "ZC706:" (checkbox), "VC707:" (checkbox), and "Board IP Addr:" (text input). An "Add" button is at the bottom of the form. Below the form, a note states: "This and other pages can be overridden and customized." At the very bottom, a footer indicates: "BuildBot (0.8.9) working for the Connectal project. Page built: Fri 05 Dec 2014 20:43:09 (GMT)".

Before submitting a project, you must sign in using your github credentials. We do not store credentials, but pass them through to github. Next, enter a name for the project, which will be used for subsequent build requests through the Buildbot web interface. The project must be in a publicly accessible git-hub repository, whose Repo location is entered beginning with git:// as follows git://github.com/connectal-examples/leds.git. If the project makefile is not in the root directory of the repository, enter its relative path in the 'Path' field of the form. If a particular branch or revision number are desired, enter these as well. Check the button to select the build target. If you have selected a zynq-based platform and would like the tool-chain to automatically program the device and execute the design as its final step, then enter the IP address of your board. This works only because adb doesn't require authentication. SSH keys required to run on PCIe-based platforms are not currently supported. Finally, don't forget to click 'Add'. If the project name has already been used, you will be prompted to enter a new one at this point.

Compiling Locally

Before compiling a project locally, you will need to install the toolchain. After setting the `CONNECTALDIR` to the root of the connectal source tree, enter the command `make`

Running the Design

When preparing a zedboard:

- Following the directions in the zynq-boot git repo, create an Android boot SD card files.
- Write the SD card files into the first partition of an SD card (FAT format).
- Verify that the boot mode jumpers on the board are JP8/9/10 (MIO3/4/5) == 110 for booting from SD card.
- Verify that the PS_MIO0 pull-down jumper on the board JP6 is shorted.
- Connect a USB cable to the 'UART' connector. Use the 'connectable' program from the connectable git repo to display linux console output.
- Connect an ethernet cable (linux uses DHCP to get an IP address during boot).

When preparing a PCIe board:

- Attach a USB cable to the JTAG port

Clocking Your Design

Every board has a main clock, which is the clock exposed by "exposeCurrentClock". It has a default value, but that value can be overridden with `CONNECTALFLAG -mainclockperiod`, which is an integer specified in nanoseconds.

There is a second clock available, which I called "derivedClock" because it was derived from the main clock. You can specify its clock period with `-derivedclockperiod`, which is a float specified in nanoseconds. I'm using an MMCM, which has one clock specified by a fractional divisor.

On PCIe-connected boards, the main clock frequency defaults to the PCIe user clock frequency (125MHz for gen1, 250MHz for gen2). But you can override that, in which case your hardware is connected to PCIe via sync FIFOs.

You are responsible for any synchronization required between the main and derived clock domains.

There are two ways to get access to the derivedClock in your design.

IMPORT_HOST_CLOCKS

This is simpler, and preserves the synthesis boundary on `mkConnectalTop`. See examples/echoslow

In the Makefile

CONNECTALFLAGS += -D IMPORT_HOST_CLOCKS

Add `”:host.derivedClock,host.derivedReset”` to `H2S_INTERFACES`:

`H2S_INTERFACES = Echo:EchoIndication:host.derivedClock,host.derivedReset`

Or just pass in `host.derivedClock` and create a reset locally.

Look at the generated `<board>/generatedbsv/Top.bsv` to see how this changed the generated code.

IMPORT_HOSTIF

This option is only useful for Zynq, in order to get access to other PS7 interfaces and clocks that are not part of the standard portal interface, e.g., I2C.

CONNECTALFLAGS += -D IMPORT_HOSTIF

Add `”:host”` to `H2S_INTERFACES`.

See `tests/test_sdio1` for an example of the use of `IMPORT_HOSTIF`, though it has a manually written `Top.bsv`.

Setting Zynq Clock Speeds

Default Clock Speeds

The default values for `-mainclockperiod` and `-derivedclockperiod` for the each board are in the JSON files in the `boardinfo` directory.

Connectal Makefile Variables

Makefile.connectal

A Connectal design imports `Makefile.connectal` into its Makefile in order to drive the build.

A number of variables are used to control the build and parameters of the design.

Environment Variables

CONNECTALDIR

Points to the location where `Makefile.connectal` and the connectal tools are installed.

Make Variables Defining the Application

BOARD

This is typically set from the suffix of the build target, e.g., `make build.zedbard` defines `BOARD=zedboard`.

INTERFACES

Specifies for which interfaces to generate `c/c++/bsv` proxies and wrappers.

NUMBER_OF_MASTERS

Number of DMA masters in the design. Defaults to 1.

PIN_TYPE

BSV interface of exported pins. Defaults to Empty. BSV type `bsv:typedef::PinType` is defined from `PIN_TYPE`.

PIN_TYPE_INCLUDE

Which BSV package to import to get the declaration of PIN_TYPE.

PINOUT_FILE

Which pin usage JSON files to pass to makefilegen.py as :option::--pinout options.

BSVFILES

Lists the BSV files to scan when processing INTERFACES.

CPPFILES

Lists the C/C++ files that implement the application.

CPPFILES2

Lists the C/C++ files that implement the (optional) second executable of an application. For example, a daemon that coordinates access to the hardware.

PORTAL_DUMP_MAP

Specifies the option to provide to pcieflat to annotate PCIe traces with portal numbers and method names. Uses generatedDesignInterfaceFile.json.

Auto Top

S2H_INTERFACES**H2S_INTERFACES****MEM_READ_INTERFACES****MEM_WRITE_INTERFACES**

Controlling the Build

CONNECTALFLAGS

Flags to pass to makefilegen.py. See _invocation_makefilegen.py for its options.

V

Controls verbosity of the build. V=1 for verbose.

USE_BUILDDCACHE

Define USE_BUILDDCACHE=1 to use buildcache. Except fpgamake seems to use buildcache anyway.

BUILDDCACHE

Location of buildcache script.

BUILDDCACHE_CACHEDIR

To specify an alternate location for the buildcache cache files.

IPDIR

Specifies into which directory to generate IP cores. This allows generated cores to be shared between designs when the FPGA part and core parameters match.

MAIN_CLOCK_PERIOD

Bound to the clock period, in nanoseconds, of the clock domain of mkConnectalTop.

Defaults to 8ns for vc707 and kc705.

Defaults to 10ns for zedboard.

Defaults to 5ns for zc706.

DEFAULT_DERIVED_CLOCK_PERIOD

Bound to the default clock period, in nanoseconds, of the derived clock provided via HostInterface to mkConnectalTop. Defaults to half the period, twice the frequency of the main clock.

DERIVED_CLOCK_PERIOD

Bound to the clock period, in nanoseconds, of the derived clock provided via HostInterface to mkConnectalTop. Defaults to DEFAULT_DERIVED_CLOCK_PERIOD.

BURST_LEN_SIZE

Controls width of fields specifying memory request burst lengths. Defaults to 8.

RUNPARAM

Specifies the name or IP address of the machine on which to run the application, e.g.:

```
make RUNPARAM=192.168.168.100 run.android
```

Top Level Make Targets

build.%

Builds software and bitfile for the specified board name, e.g.,:

```
make build.zedboard
```

run.%

Programs the FPGA and runs the application using the build for the specified board name. Uses :make:var:RUNPARAM. For example,:

```
make RUNPARAM=sjl0 run.vc707
```

Intermediate Make Targets

verilog

Runs the build up through generation of verilog from BSV. Requires BOARD to be defined.

bits

Generates the FPGA bit file from the design. Requires BOARD to be defined.

bsim

For BOARD=bluesim, generates the simulation executable.

xsim

For BOARD=xsim, generates the simulation executable.

android.exe

Builds the software executable for boards using Android.

ubuntu.exe

Builds the software executable for boards using Ubuntu/CentOS.

bsim_exe

Builds the software executable for bluesim.

gentarget

This step creates the board directory and Makefile.

prebuild

Additional steps needed before making verilog, etc. Use this target for dependences such as constraint file and IP core generation that need to be run before the design is built. This is a :: dependence, so you can specify it multiple times.

Makefile.connectal.build

The target Makefile, which generated by makefilegen.py (which is invoked by Makefile.connectal), includes Makefile.connectal.build.

Make targets

bits

Builds the FPGA bitstream.

exe

Builds the application software (executable or shared library).

run

Programs the FPGA and Runs the application.

Board-specific targets

ubuntu.exe

Builds the application for boards using Ubuntu/Debian/CentOS Linux. Also used with xsim.

android.exe

Builds the application for boards using Android (e.g., zedboard).

bsim_exe

Builds the application for use with bluesim

bsim

Builds the hardware for Bluespec bluesim simulator

xsim

Builds the hardware for Xilinx xsim simulator

Environment Variables

RUNPARAM

Specifies the name or address of the host on which to run the application.

NOFPGAJTAG

(Deprecated, use :envvar::*NOPROGRAM*. See also *noprogram*.) If NOFPGAJTAG is defined, then fpgajtag is not called by ubuntu.exe or android.exe.

NOPROGRAM

If NOPROGRAM is defined, then the FPGA is not automatically programmed by ubuntu.exe or android.exe. See also *noprogram*.

Variables Controlling the Build

V

Binding V=1 makes the build verbose, displaying the commands that are run while building the project.

CONNECTAL_SHARED

Bind to 1 to build a shared library instead of an executable application.

CONNECTAL_NOHARDWARE

Binding to 1 suppresses creation of FPGA bitstream.

CONNECTAL_NDK_PARAM

Specifies options to pass to ndk-build.

Variables Controlling the Bluespec Compiler

BSCOPTFLAGS

Specifies bsc optimization flags.

BSPATHFLAGS

Specifies directories used to store output files from Bluespec compilation.

BSCFLAGS_COMMON

Specifies Connectal-specific bsc flags.

BSCFLAGS_PROJECT

Specifies project-specific bsc flags, -p argument passed to makefilegen.py via CONNECTALFLAGS.

BSCFLAGS_EXTRA

Specifies additional flags to pass to bsc.

Bluesim-related Variables

BSC_CXXFLAGS

CXXFLAGS passed to C++ compiler when building application for bluesim target

Vivado-related Variables

VIVADOFLAGS

Specifies options to pass when running vivado.

Buildcache related Variables

USE_BUILDCACHE

Bind to 1 to use buildcache.

Variables Controlling the Application

CONNECTAL_DEBUG

Bind CONNECTAL_DEBUG=1 to run the application under gdb.

RUN_ARGS

Specifies arguments to pass to the application when invoking it.

Connectal Tools

Invocation of makefilegen.py

The **makefilegen.py** script generates project Makefiles. Normally, it is invoked from Makefile.connectal and passed options via *CONNECTALFLAGS*.

Project Options

- B board, --board=board**
Specifies which board to build for.
- O os, --option=os**
Specifies which operating system to support. Usually derived from :option:board.
- project-dir=dir**
Specifies the directory in which to creat the Makefile that performs the build. Creates the directory if it does not exist.

Source Options

bsvfile

This positional argument specifies which BSV files to parse for type and interface declarations.

- s foo.cpp, --source=foo.cpp**
Specifies C and C++ files to include in the application.
- ipdir=dir**
Specifies where to store IP cores generated by Vivado or Quartus
- cachedir=dirname**
Specifies directory to be used by *buildcache*
- nocache**
Disables *buildcache*
- D var=value**
Binds var to value in software, BSV, Verilog, and Tcl contexts.
- v, --verbose**
Verbose operation

Interface Options

- interfaces=...**
To be written...
- dump_map**
Generate a JSON file describing the portals, for use in tracing tools such as *pcieflat*.

Software Options

- cflags=flags**
Specifies build flags for the C/C++ compiler
- I dir -cinclude=dir**
- l lib, --clib=lib**
Link the application with library lib.
- S libfile, --clibfiles=libfile.a**
Link the application with library file *libfile.a*.
- L libdir, --clibdir=libdir**
Addes *libdir* to the software library path.

--nonstrict

Pass `-Wall` to gcc instead of `-Werror`.

--shared

Specifies to build a shared library instead of an executable.

--nohardware

Suppresses build of hardware, for software-only projects.

--stl=stltype

Specifies which C++ Standard Template Library to use for Android. The choices are:

- `stlport_static`: STLport runtime C++ exceptions and RTTI and Standard Template Library
- `stlport_shared`
- `gnustl_static`: GNU STL runtime C++ exceptions and RTTI and Standard Template Library
- `gnustl_shared`
- `c++_static`: LLVM libc++ runtime C++ exceptions and RTTI and Standard Template Library
- `c++_shared`
- `gabi++_static`: GAbi++ runtime C++ exceptions and RTTI
- `gabi++_shared`

Hardware Options

--bsvpath=dirname

Adds *dirname* to the BSV module import path.

-b options, --bscflags=options

Flags to pass to the BSV compiler.

-V file.v, --verilog=file.v

Specifies verilog file to include in the design

--pinfo=project.json

This does what?

--pinout=pinusage.json

Specifies connection of pins on the board to top level signals of the design. A board-specific constraint file will be generated from pinusage.json and the boardinfo json file, and will be added to constraint and implconstraint.

See `:make:var::PINOUT_FILE`, whose values are passed as `-pinout` options to `makefilegen`.

--constraint=file.xdc

Specifies synthesis phase constraint file.

--implconstraint=file.xdc

Specifies implementation (place and route) constraint file.

-P modulename, --partition=modulename

Directs *fpgamake* to build a separate netlist for *modulename*. With *buildcache*, Reduces build times if module changes infrequently.

Xilinx Options

--xci=core.xci

Specifies IP core to include in the design. (Xilinx only)

Altera Options

- qip=core.qip**
Specifies IP core to include in the design. (Altera only)
- qsf=settings.qsf**
Specifies Altera Quartus settings.

Partial Reconfiguration Options

- prtop=mkTop.dcp**
Specifies filename of previously built top level.
- prvariant=name**
...
- reconfig=modulenames**
...

Bluesim Options

- q, --qtused**
Link the bluesim *bsim* executable with libQt.
- m foo.cpp, --bsimsource foo.cpp**
Specifies additional sources to compile into the *bsim* executable. If you are using Bluespec import “BDPI” or SystemVerilog “DPI”/“DPI-C”, you will need to link additional sources into the simulator.

Note: These files are currently compiled with g++, even if they are C files. You will need to use extern “C” to export symbols to the simulator.

Xsim Options

- xelabflags=flags**
Options to pass to *xelab*
- xsimflags**
Options to pass to *xsim*

Clocking Options

- mainclockperiod**
Specifies the period, in nanoseconds, of the main clock. Must be an integer.

On Zynq boards,

On PCIe-connected boards, if the main clock period differs from the PCIe clock period, then the design’s portals and DMA ports will automatically be connected via SyncFIFO’s.

Each boardinfo JSON file specifies the default value for mainclockperiod.
- derivedclockperiod**
Connectal also makes a second clock available, host.derivedClock. The period of this clock is

Each boardinfo JSON file specifies the default value for mainclockperiod.

Invocation of generate-constraints.py

```
usage: generate-constraints [-h] [--boardfile BOARDFILE]
                             [--pinoutfile PINOUTFILE] [-b BIND] [-o OUTPUT]
                             [-f FPGA]
```

Options:

--boardfile Board description file (json)
--pinoutfile=[] Project description file (json)
-b=[], --bind=[] Bind signal group to pin group
-o, --output Write output to file
-f=xilinx, --fpga=xilinx Target FPGA Vendor

pcieflat

Usage:

```
pcieflat
```

Dumps trace of PCIE transactions.

The trace contains the following columns:

col	description
dirtytype	TLP direction and type
timestamp	Timestamp of TLP, in main clock cycles
delta	Number of cycles since previous TLP
type	type:pktype:format
be	Byte enables
hit	BAR hit
eof	End of frame
sof	Start of frame
address	Request physical address
tag	Request tag

The first column, *dirtytype*, contains RX if received by FPGA, TX if transmitted by FPGA, qq for start of request, pp for start of response, and cc for continuation.

The second column, *timestamp*, displays the timestamp in terms of a 64-bit counter running at the PCIe user clock frequency (125MHz for gen1, 250MHz for gen2 and gen3).

The third column, *delta*, displays the number of cycles since the previous TLP. If a TLP is transmitted and received on the same cycle, then the transmitted TLP will have a delta of 0. The first TLP shown in the trace will have a delta of 0.

The fourth column shows the type of the TLP. If it is the start of a transaction, it will be of the form *type:tlppktype:tlpformat*, where the types are:

- CpuRReq: read request from CPU to FPGA
- CpuWReq: write request from CPU to FPGA
- CpuRRsp: read response from FPGA to CPU
- DmaRReq: read request from FPGA to CPU DRAM
- DmaWReq: write request from FPGA to CPU DRAM

- DmaRRsp: read response from CPU DRAM to FPGA
- CpuRCon: continuation data sent from FPGA to CPU (continuation of CpuRRsp or DmaWReq)
- DmaRCon: continuation data sent from CPU DRAM to FPGA (continuation of DmaRRsp or CpuWReq)
- Interru: interrupt message from FPGA to CPU

The TLP *pkttype* is one of the following:

- MRW: Memory read/write
- COMP: additional data (“completion”) of transaction

The TLP *format* is one of the following:

- MEM_WRITE_3DW_DATA: 96-bit write request header containing 32-bit address and 32-bit data
- MEM_WRITE_4DW_DATA: 128-bit write request header containing 64-bit address and no data
- MEM_READ_3DW: 96-bit read request header containing 32-bit address
- MEM_READ_4DW: 128-bit read request header containing 64-bit address

The *sof* and *eof* flags indicate the start and end TLPs of each transaction.

Diagnosing Crashes

If the machine crashed, look for transactions that were started but not ended.

These will generally fall before very large deltas, where the machine was rebooted before it had any more interactions with the FPGA.

Example Trace

Here are some excerpts from the output of *pcieflat* after running *tests/memserver_copy*:

pcieflat: devices are ['/dev/portal_0_6', '/dev/portal_0_5', '/dev/portal_0_4', '/dev/portal_0_3', '/dev/portal_0_2', '/dev/portal_0_1', '/dev/portal_0_0']									
	ts	delta	response	XXX	tlp	address	off	be	
			pkttype	format	foo	(be	hit	eof	sof)
RXpp	513858932		0	DmaRRsp:COMP:MEM_WRITE_3DW_DATA	09 0x4	tlp(ffff	0	0	1)
TXcc	513858932		0	CpuRCon	10 0x8	tlp(ffff	0	0	0)
RXcc	513858933		1	DmaRCon	08 0x4	tlp(ffff	0	0	0)
TXcc	513858933		0	CpuRCon	10 0x8	tlp(ffff	0	1	0)
RXcc	513858934		1	DmaRCon	08 0x4	tlp(ffff	0	0	0)
RXcc	513858935		1	DmaRCon	08 0x4	tlp(ffff	0	0	0)
RXcc	513858936		1	DmaRCon	08 0x4	tlp(ffff	0	1	0)
RXpp	513858938		2	DmaRRsp:COMP:MEM_WRITE_3DW_DATA	09 0x4	tlp(ffff	0	0	1)
RXcc	513858939		1	DmaRCon	08 0x4	tlp(ffff	0	0	0)
RXcc	513858940		1	DmaRCon	08 0x4	tlp(ffff	0	0	0)
RXcc	513858941		1	DmaRCon	08 0x4	tlp(ffff	0	0	0)
RXcc	513858942		1	DmaRCon	08 0x4	tlp(ffff	0	1	0)
TXqq	513858949		7	DmaWReq: MRW:MEM_WRITE_4DW_DATA	11 0x8	tlp(ffff	0	0	1)
TXcc	513858950		1	CpuRCon	10 0x8	tlp(ffff	0	0	0)
TXcc	513858951		1	CpuRCon	10 0x8	tlp(ffff	0	0	0)
TXcc	513858952		1	CpuRCon	10 0x8	tlp(ffff	0	0	0)
TXcc	513858953		1	CpuRCon	10 0x8	tlp(ffff	0	0	0)
TXcc	513858954		1	CpuRCon	10 0x8	tlp(ffff	0	0	0)
TXcc	513858955		1	CpuRCon	10 0x8	tlp(ffff	0	0	0)
TXcc	513858956		1	CpuRCon	10 0x8	tlp(ffff	0	0	0)
TXcc	513858957		1	CpuRCon	10 0x8	tlp(ffff	0	1	0)

It may contain DMA Read Requests:

TXqq	513859009	1	DmaRReq: MRW:MEM_READ__4DW	11	0x8	tlp(ffff 0 1 1)	address: 000000039
TXqq	513859010	1	DmaRReq: MRW:MEM_READ__4DW	11	0x8	tlp(ffff 0 1 1)	address: 000000039

The trace may contain interrupts:

TXqq	513865009	12	Interru: MRW:MEM_WRITE_3DW_DATA	11	0x8	tlp(ffff 0 1 1)	fee00000 0 be (1
------	-----------	----	---------------------------------	----	-----	-----------------	------------------

It will also contain reads and writes from the CPU to the FPGA:

RXqq	513874994	9985	CpuRReq: MRW:MEM_READ__3DW	09	0x4	tlp(ffff 4 1 1)	df400000 0 be (1
TXpp	513875008	14	CpuRRsp:COMP:MEM_WRITE_3DW_DATA	11	0x8	tlp(ffff 4 1 1)	
RXqq	513875159	151	CpuRReq: MRW:MEM_READ__3DW	09	0x4	tlp(ffff 4 1 1)	df402000 800 be (1
TXpp	513875173	14	CpuRRsp:COMP:MEM_WRITE_3DW_DATA	11	0x8	tlp(ffff 4 1 1)	
RXqq	513875323	150	CpuRReq: MRW:MEM_READ__3DW	09	0x4	tlp(ffff 4 1 1)	df401000 400 be (1
TXpp	513875337	14	CpuRRsp:COMP:MEM_WRITE_3DW_DATA	11	0x8	tlp(ffff 4 1 1)	
RXqq	513876236	899	CpuRReq: MRW:MEM_READ__3DW	09	0x4	tlp(ffff 4 1 1)	df40000c 3 be (1
TXpp	513876250	14	CpuRRsp:COMP:MEM_WRITE_3DW_DATA	11	0x8	tlp(ffff 4 1 1)	
RXqq	513876449	199	CpuRReq: MRW:MEM_READ__3DW	09	0x4	tlp(ffff 4 1 1)	df400020 8 be (1
TXpp	513876463	14	CpuRRsp:COMP:MEM_WRITE_3DW_DATA	11	0x8	tlp(ffff 4 1 1)	
RXqq	513883818	7355	CpuRReq: MRW:MEM_READ__3DW	09	0x4	tlp(ffff 4 1 1)	df40000c 3 be (1
TXpp	513883832	14	CpuRRsp:COMP:MEM_WRITE_3DW_DATA	11	0x8	tlp(ffff 4 1 1)	
RXqq	513883976	144	CpuWReq: MRW:MEM_WRITE_3DW_DATA	09	0x4	tlp(ffff 4 1 1)	df400004 1 be (1
RXqq	513884007	31	CpuRReq: MRW:MEM_READ__3DW	09	0x4	tlp(ffff 4 1 1)	df40200c 803 be (1
TXpp	513884021	14	CpuRRsp:COMP:MEM_WRITE_3DW_DATA	11	0x8	tlp(ffff 4 1 1)	
RXqq	513884165	144	CpuWReq: MRW:MEM_WRITE_3DW_DATA	09	0x4	tlp(ffff 4 1 1)	df402004 801 be (1

The trace will end with a summary of the kinds of PCIe transactions:

```
{ 'DmaWReq': 115, 'CpuRReq': 45, 'Interru': 3, 'CpuRRsp': 45, 'CpuRCon': 922, 'CpuWReq': 15, 'DmaRReq': 2302 }
```

Invocation of topgen.py

```
usage: topgen [-h] [--project-dir PROJECT_DIR] [--filename FILENAME]
              [--topname TOPNAME] [--ifcnames IFCNAMES] [--pintype PINTYPE]
              [--interface INTERFACE] [--portalclock PORTALCLOCK]
              [--importfiles IMPORTFILES] [--portname PORTNAME]
              [--wrapper WRAPPER] [--proxy PROXY] [--memread MEMREAD]
              [--memwrite MEMWRITE] [--cnoc] [--integratedIndication]
```

Options:

--project-dir project directory

--filename=Top.bsv name of generated file

--topname=mkConnectalTop name of generated module

--ifcnames=IfcNames name of interface names enum type and file

--pintype=[] Type of pins interface

--interface=[] exported interface declaration

--portalclock Portal clock source

--importfiles=[] added imports

--portname=[] added portal names to enum list


```

--wrapper=[]      exported wrapper interfaces
--proxy=[]        exported proxy interfaces
--memread=[]      memory read interfaces
--memwrite=[]     memory read interfaces
--cnoc=False      generate mkCnocTop
--integratedIndication=False  indication pipes instantiated in user module

```

Connectal BSV Libraries

Address Generator

One of the common patterns that leads to long critical paths in designs on the FPGA are counters and comparisons against counters. This package contains a module for generating the sequence of addresses used by a memory read or write burst, along with a field indicating the last beat of the burst.

```
struct AddressGenerator :: AddrBeat#(numeric type addrWidth)
```

```
addr → Bit#(addrWidth)
    The address for this beat of the request.
```

```
bc → Bit#(BurstLenSize)
```

```
tag → Bit#(MemTagSize)
```

```
last → Bool
```

```
interface AddressGenerator :: AddressGenerator (numeric type addrWidth, numeric type dataWidth)
```

```
request → Put#(PhysMemRequest#(addrWidth))
    The interface for requesting a sequence of addresses.
```

```
addrBeat → Get#(AddrBeat#(addrWidth))
    The interface for getting the address beats of the burst. There is one pipeline cycle from the request to the first address beat.
```

```
module AddressGenerator :: mkAddressGenerator → (AddressGenerator#(addrWidth, dataWidth))
    Instantiates an address generator.
```

Arith Package

The Arith package implements some functions that correspond to infix operators.

```
function Arith :: booland (Bool x1, Bool x2) → Bool
    Returns logical “and” of inputs. Named to avoid conflict with the Verilog keyword “and”.
```

```
function Arith :: boolor (Bool x1, Bool x2) → Bool
    Returns logical “or” of inputs. Named to avoid conflict with the Verilog keyword “or”.
```

```
function Arith :: eq (a x1, a x2) → Bool
```

```
function Arith :: add (a x1, a x2) → a
    Returns sum of inputs. Requires Arith#(a).
```

function Arith: **mul** (*a x1, a x2*) → a
Returns product of inputs. Requires Arith#(a).

function Arith: **rshift** (*Bit#(b) x1, Integer i*) → Bit#(b)
Returns input right shifted by i bits.

function Arith: **a) vadd** (*Vector#(n, a) x1, Vector#(n, a) x2*) → Vector#(n,
Returns sum of input vectors.

function Arith: **a) vmul** (*Vector#(n, a) x1, Vector#(n, a) x2*) → Vector#(n,
Returns element-wise product of input vectors.

function Arith: **Bit# (b)) vrshift**(*Vector#(n, Bit#(b)) x1, Integer i*) → Vector#(n,
Right shifts the elements of the input vector by i bits.

AxiStream Package

AXI Stream Interfaces

interface AxiStream: **AxiStreamMaster** (*numeric type dataWidth*)
AXI stream source with dataWidth data bits.

tdata → Bit#(dsz)
Returns the data from this beat if tvalid is asserted, otherwise returns undefined.

tkeep → Bit#(TDiv#(dsz,8))
Returns the byte enables from this beat if tvalid is asserted, otherwise returns undefined.

tlast → Bit#(1)
Indicates if this is the last data beat of this transaction tvalid is asserted, otherwise returns undefined.

tready (*Bit#(1) v*) → Action
When tvalid and tready are both asserted the current data is consumed. The value passed to tready may not depend on the output of tvalid.

tvalid → Bit#(1)
Asserted when the data is valid.

interface AxiStream: **AxiStreamSlave** (*numeric type dataWidth*)
AXI stream sink with dataWidth data bits.

tdata (*Bit#(dsz) data*) → Action
The data passed from the source if tvalid is asserted, otherwise undefined..

tkeep (*Bit#(TDiv#(dsz, 8)) keep*) → Action
The byte enables passed from the source if tvalid is asserted, otherwise undefined.

tlast (*Bit#(1) last*) → Action
Indicates if this is the last data beat of this transaction tvalid is asserted, otherwise returns undefined.

tready → Bit#(1)
Return 1 if ready to receive data, 0 otherwise.

When tvalid and tready are both asserted the current data is consumed. The value passed to tready may not depend on the output of tvalid.

tvalid (*Bit#(1) v*) → Action
Indicates the data from the source is valid.

Connectable Type Instances

```
instance AxiStream: :Connectable (AxiStreamMaster#(dataWidth), AxiStreamSlave#(dataWidth))
instance AxiStream: :ToGetM (AxiStreamMaster#(asz), Bit#(asz))
instance AxiStream: :ToPutM (AxiStreamSlave#(asz), Bit#(asz))
```

AXI Stream Type Classes and Instances

```
typeclass AxiStream: :ToAxiStream (type atype, type btype)
typeclass AxiStream: :... bsv: function: : atype toAxiStream (btype b)
```

Convert to an AXI stream interface.

```
typeclass AxiStream: :MkAxiStream (type atype, type btype)
```

```
module mkAxiStream (btype b) → (atype)
  Create a module with an AXI Stream interface.
```

```
instance AxiStream: :MkAxiStream (AxiStreamMaster#(dsize), FIFO#(Bit#(dsize)))
```

```
module mkAxiStream (FIFO#(Bit#(dsize)) f) → (AxiStreamMaster#(dsize))
  Create an AXI Stream master from a FIFO of bits
```

```
instance AxiStream: :MkAxiStream (AxiStreamSlave#(dsize), FIFO#(Bit#(dsize)))
```

```
module mkAxiStream (FIFO#(Bit#(dsize)) f) → (AxiStreamSlave#(dsize))
  Create an AXI Stream slave from a FIFO of bits
```

```
instance AxiStream: :MkAxiStream (AxiStreamMaster#(dsize), FIFO#(Bit#(dsize)))
```

```
module mkAxiStream (FIFO#(Bit#(dsize)) f) → (AxiStreamMaster#(dsize))
  Create an AXI Stream master from a FIFO of MemDataF
```

```
instance AxiStream: :MkAxiStream (AxiStreamSlave#(dsize), FIFO#(MemDataF#(dsize)))
```

```
module mkAxiStream (FIFO#(MemDataF#(dsize)) f) → (AxiStreamSlave#(dsize))
  Create an AXI Stream slave from a FIFO of MemDataF
```

CtrlMux Package

```
module CtrlMux: :mkInterruptMux (Vector#(numPortals, MemPortal#(aw, dataWidth)) portals) →
  (ReadOnly#(Bool))
```

Used by BsimTop, PcieTop, and ZynqTop. Takes a vector of MemPortals and returns a boolean indicating whether any of the portals has indication method data available.

```
module CtrlMux: :mkSlaveMux (Vector#(numPortals, MemPortal#(aw, dataWidth)) portals) → (Phys-
  MemSlave#(addrWidth, dataWidth))
```

Takes a vector of MemPortals and returns a PhysMemSlave combining them.

HostInterface Package

The HostInterface package provides host-specific typedefs and interfaces.

Host-Specific Constants

typedef HostInterface::DataBusWidth

Width in bits of the data bus connected to host shared memory.

typedef HostInterface::PhysAddrWidth

Width in bits of physical addresses on the data bus connected to host shared memory.

typedef HostInterface::NumberOfMasters

Number of memory interfaces used for connecting to host shared memory.

Host-Specific Interfaces

interface HostInterface::BsimHost

Host interface for the bluesim platform

interface HostInterface::PcieHost

Host interface for PCIe-attached FPGAs such as vc707 and kc705

interface HostInterface::ZynqHost

Host interface for Zynq FPGAs such as zedboard, zc702, zc706, and zybo.

The Zc706 is a ZynqHost even when it is plugged into a PCIe slot.

Application-Specific Types

typedef HostInterface::PinType

Specifies the type of the application pins interface. Defined from PIN_TYPE.

Leds Package

interface Leds::LEDS

typedef Leds::LedsWidth

Defined to be the number of default LEDs on the FPGA board.

The Zedboard has 8, Zc706 has 4, ...

Leds::leds → Bit#(LedsWidth)

MemPortal Package

mkMemPortal Module

module MemPortal::mkMemPortal (*Bit#(slaveDataWidth) ifcId, PipePortal#(numRequests, numIndications, slaveDataWidth) portal*) → (MemPortal#(slaveAddrWidth, slaveDataWidth))

Takes an interface identifier and a PipePortal and returns a MemPortal.

MemReadEngine Package

module MemReadEngine :: **mkMemReadEngine** (MemReadEngine (busWidth, userWidth, cmdQDepth, numServers)

Creates a MemReadEngine with default 256 bytes of buffer per server.

module MemReadEngine :: **mkMemReadEngineBuff** (Integer bufferSizeBytes) → (MemReadEngine#(busWidth, userWidth, cmdQDepth, numServers)

Creates a MemReadEngine with the specified buffer size.

MemTypes Package

Constants

typedef MemTypes :: **Bit#(32) SGLId**

typedef MemTypes :: **44 MemOffsetSize**

typedef MemTypes :: **6 MemTagSize**

typedef MemTypes :: **8 BurstLenSize**

typedef MemTypes :: **32 MemServerTags**

typedef MemTypes :: **TDiv#(DataBusSize, 8) ByteEnableSize**

Data Types

struct MemTypes :: **PhysMemRequest#(numeric type addrWidth, dataWidth)**

A memory request containing a physical memory address

addr → Bit#(addrWidth)

Physical address to read or write

burstLen → Bit#(BurstLenSize)

Length of read or write burst, in bytes. The number of beats of the request will be the burst length divided by the physical width of the memory interface.

tag → Bit#(MemTagSize)

firstbe → Bit#(TDiv#(dataWidth, 8))

lastbe → Bit#(TDiv#(dataWidth, 8))

If BYTE_ENABLES is defined as aBSV preprocessor macro, byte write enables are added to PhysMemRequest, into two fields: firstbe and lastbe. The idea is to enable writing any number of contiguous bytes even if it is less than the width of the shared memory data bus.

These have roughly the same semantics as in PCIE. The write enable in firstbe apply to the first beat of a burst request and those in lastbe apply to the last beat of a multi-beat burst request. Intervening beats of a burst request enable the write of all beats of that burst.

struct MemTypes :: **MemRequest**

A logical memory read or write request. The linear offset of the request will be translated by an MMU according to the specified scatter-gather list.

sglId → SGLId

Indicates which scatter-gather list the MMU should use when translating the address

offset → Bit#(MemOffsetSize)

Linear byte offset to read or write.

burstLen → Bit#(BurstLenSize)

Length of read or write burst, in bytes. The number of beats of the request will be the burst length divided by the physical width of the memory interface.

tag → Bit#(MemTagSize)

firstbe → Bit#(ByteEnableSize)

lastbe → Bit#(ByteEnableSize)

If `BYTE_ENABLES` is defined as an ABSV preprocessor macro, byte write enables are added to `PhysMemRequest`, into two fields: `firstbe` and `lastbe`. The idea is to enable writing any number of contiguous bytes even if it is less than the width of the shared memory data bus.

These have roughly the same semantics as in PCIE. The write enable in `firstbe` apply to the first beat of a burst request and those in `lastbe` apply to the last beat of a multi-beat burst request. Intervening beats of a burst request enable the write of all beats of that burst.

struct MemTypes::MemData#(numeric type dsz)

One beat of the payload of a physical or logical memory read or write request.

data → Bit#(dsz)

One data beat worth of data.

tag → Bit#(MemTagSize)

Indicates to which request this beat belongs.

last → Bool

Indicates that this is the last beat of a burst.

struct MemTypes::MemDataF#(numeric type dsz)

One beat of the payload of a physical or logical memory read or write request. Used by `MemReadEngine` and `MemWriteEngine`.

data → Bit#(dsz)

One data beat worth of data.

tag → Bit#(MemTagSize)

Indicates to which request this beat belongs.

first → Bool

Indicates that this is the first data beat of a request.

last → Bool

Indicates that this is the last data beat of a request.

Physical Memory Clients and Servers

interface MemTypes::PhysMemMaster (numeric type addrWidth, numeric type dataWidth)

The physical memory interface exposed by `MemMaster`. For example, connects via AXI to Zynq or via PCIe to x86 memory.

read_client → PhysMemReadClient#(addrWidth, dataWidth)

write_client → PhysMemWriteClient#(addrWidth, dataWidth)

interface MemTypes::PhysMemReadClient (numeric type asz, numeric type dsz)

readReq → Get#(PhysMemRequest#(asz))

```
readData → Put#(MemData#(dsz))
```

```
interface MemTypes : PhysMemWriteClient (numeric type asz, numeric type dsz)
```

```
writeReq → Get#(PhysMemRequest#(asz))
```

```
writeData → Get#(MemData#(dsz))
```

```
writeDone → Put#(Bit#(MemTagSize))
```

```
interface MemTypes : PhysMemSlave (numeric type addrWidth, numeric type dataWidth)
```

```
read_server → PhysMemReadServer#(addrWidth, dataWidth)
```

```
write_server → PhysMemWriteServer#(addrWidth, dataWidth)
```

```
interface MemTypes : PhysMemReadServer (numeric type asz, numeric type dsz)
```

```
readReq → Put#(PhysMemRequest#(asz))
```

```
readData → Get#(MemData#(dsz))
```

```
interface MemTypes : PhysMemWriteServer (numeric type asz, numeric type dsz)
```

```
writeReq → Put#(PhysMemRequest#(asz))
```

```
writeData → Put#(MemData#(dsz))
```

```
writeDone → Get#(Bit#(MemTagSize))
```

Memory Clients and Servers

These clients and servers operate on logical addresses. These are translated by an MMU before being issued to system memory.

```
interface MemTypes : MemReadClient (numeric type dsz)
```

The system memory read interface exported by a client of MemServer, such as MemReadEngine.

```
readReq → Get#(MemRequest)
```

```
readData → Put#(MemData#(dsz))
```

```
interface MemTypes : MemWriteClient (numeric type dsz)
```

The system memory write interface exported by a client of MemServer, such as MemWriteEngine.

```
writeReq → Get#(MemRequest)
```

```
writeData → Get#(MemData#(dsz))
```

```
writeDone → Put#(Bit#(MemTagSize))
```

```
interface MemTypes : MemReadServer (numeric type dsz)
```

The system memory read interface exported by MemServer.

```
readReq → Put#(MemRequest)
```

```
readData → Get#(MemData#(dsz))
```

```
interface MemTypes : MemWriteServer (numeric type dsz)
```

The system memory write interface exported by MemServer.

```
writeReq → Put#(MemRequest)
```

writeData → Put#(MemData#(dsz))

writeDone → Get#(Bit#(MemTagSize))

Memory Engine Types

struct MemTypes : : **MemengineCmd**

A read or write request for a MemReadEngine or a MemWriteEngine. MemRead and MemWrite engines will issue one or more burst requests to satisfy the overall length of the request.

sglId → SGLId

Which memory object identifier (scatter gather list ID) the MMU should use to translate the addresses

base → Bit#(MemOffsetSize)

Logical base address of the request, as a byte offset

burstLen → Bit#(BurstLenSize)

Maximum burst length, in bytes.

len → Bit#(32)

Number of bytes to transfer. Must be a multiple of the data bus width.

tag → Bit#(MemTagSize)

Identifier for this request.

Memory Engine Interfaces

interface MemTypes : : **MemWriteEngineServer** (*numeric type userWidth*)

The interface used by one client of a MemWriteEngine.

request → Put#(MemengineCmd)

done → Get#(Bool)

data → PipeIn#(Bit#(userWidth))

interface MemTypes : : **MemWriteEngine** (*numeric type busWidth, numeric type userWidth, numeric type cmdQDepth, numeric type numServers*)

A multi-client component that supports multi-burst writes to system memory.

dmaClient → MemWriteClient#(busWidth)

writeServers → Vector#(numServers, MemWriteEngineServer#(userWidth))

interface MemTypes : : **MemReadEngineServer** (*numeric type userWidth*)

The interface used by one client of a MemReadEngine.

request → Put#(MemengineCmd)

data → PipeOut#(MemDataF#(userWidth))

interface MemTypes : : **MemReadEngine** (*numeric type busWidth, numeric type userWidth, numeric type cmdQDepth, numeric type numServers*)

A multi-client component that supports multi-burst reads from system memory.

dmaClient → MemReadClient#(busWidth)

readServers → Vector#(numServers, MemReadEngineServer#(userWidth))

Memory Traffic Interfaces

interface MemTypes : : DmaDbg

getMemoryTraffic → ActionValue#(Bit#(64))

dbg → ActionValue#(DmaDbgRec)

Connectable Instances

instance MemTypes : : Connectable (MemReadClient#(dsz), MemReadServer#(dsz))

instance MemTypes : : Connectable (MemWriteClient#(dsz), MemWriteServer#(dsz))

instance MemTypes : : Connectable (PhysMemMaster#(addrWidth, busWidth), PhysMemSlave#(addrWidth, busWidth))

instance MemTypes : : Connectable (PhysMemMaster#(32, busWidth), PhysMemSlave#(40, busWidth))

MMU Package

typedef MMU : : 32 MaxNumSGLists

typedef MMU : : Bit# (TLog# (MaxNumSGLists)) SGListId

typedef MMU : : 12 SGListPageShift0

typedef MMU : : 16 SGListPageShift4

typedef MMU : : 20 SGListPageShift8

typedef MMU : : 24 SGListPageShift12

typedef MMU : : Bit# (TLog# (MaxNumSGLists)) RegionsIdx

typedef MMU : : 8 IndexWidth

Address Translation

struct MMU : : AddrTransRequest

Address translation request type

id → SGListId

Which object identifier to use.

off → Bit#(MemOffsetSize)

The address to translate.

interface MMU : : MMU (numeric type addrWidth)

An address translator

request → MMURequest

The interface of the MMU that is exposed to software as a portal.

addr → Vector#(2, Server#(AddrTransRequest, Bit#(addrWidth)))

The address translation servers

interface MMU : **MMURequest**;

The Connectal MMU maps linear offsets on objects identified by `sglId` to `dmaAddress`. It is constructed from a list of segments, where the segments are sorted by size in descending order. Each segment must be one of the supported sizes.

sglist (*Bit#(32) sglId, Bit#(32) segmentIndex, Bit#(64) addr, Bit#(32) len*) → Action

Updates the address of the segment number `segmentIndex` for object identified by `sglId`. The address has been preshifted so that the final address may be constructed by concatenating `addr` and offset within the segment.

region (*Bit#(32) sglId, Bit#(64) barr12, Bit#(32) idxOffset12, Bit#(64) barr8, Bit#(32) idxOffset8, Bit#(64) barr4, Bit#(32) idxOffset4, Bit#(64) barr0, Bit#(32) idxOffset0*) → Action

Updates the boundaries between the segments of different sizes for the object identified by `sglId`.

For example, if an offset to be translated is less than `barr12`, then the target segment is of size `SGList-PageShift12` (2^{24} bytes). If the offset is less than `barr12`, then `idxOffset12` points to the first translation table entry for segments of that size

```
pbase = offset >> segAddrSize + idxOffset segNumber = pbase + idxOffset dmaBase
= translationTable[sglId,segNumber] dmaAddress = {dmaBase[physAddrSize-segAddrSize-1:0],offset[segAddrSize-1:0]}
```

idRequest (*SpecialTypeForSendingFd fd*) → Action

Requests a new object identifier.

idReturn (*Bit#(32) sglId*) → Action

Indicates that the designated object is no longer in use. The MMU clears the translation entries for this object.

setInterface (*Bit#(32) interfaceId, Bit#(32) sglId*) → Action

This method is only implemented in software responders.

interface MMU : **MMUIndication**;

interface MMU : ... **bsv.method** : **Action idResponse** (**Bit** ((32) *sglId*);)

Response from `idRequest` indicating the new object identifier `sglId`.

configResp (*Bit#(32) sglId*) → Action

error (*Bit#(32) code, Bit#(32) sglId, Bit#(64) offset, Bit#(64) extra*) → Action

Sent from the MMU when there is a translation error.

struct MMU : **DmaErrorType**

DmaErrorNone

Code 0 indicates no error.

DmaErrorSGLIdOutOfRange_r

Code 1 indicates object identifier was out of range during a read request.

DmaErrorSGLIdOutOfRange_w

Code 2 indicates object identifier was out of range during a read request.

DmaErrorMMUOutOfRange_r

Code 3 indicates MMU identifier was out of range during a read request.

DmaErrorMMUOutOfRange_w

Code 4 indicates MMU identifier was out of range during a read request.

DmaErrorOffsetOutOfRange

Code 5 indicates offset was out of range for the designated object.

DmaErrorSGLIdInvalid

Code 6 indicates the object identifier was out of range.

DmaErrorTileTagOutOfRange

Code 7 indicates the tag was out of range for the requesting platform application tile.

```
module MMU : mkMMU (Integer iid, Bool hostMapped, MMUIndication mmuIndication) →  
                (MMU#(addrWidth))
```

Instantiates an address translator that stores a scatter-gather list to define the logical to physical address mapping.

Parameter iid is the portal identifier of the MMURequest interface.

Parameter hostMapped is true for simulation.

```
interface MMU : MemServerRequest;
```

```
  addrTrans (Bit#(32) sglId, Bit#(32) offset) → Action
```

Requests an address translation

```
  setTileState (TileControl tc) → Action
```

Changes tile status

```
  stateDbg (ChannelType rc) → Action
```

Requests debug info for the specified channel type

```
  memoryTraffic (ChannelType rc) → Action
```

```
interface MMU : MemServerIndication;
```

```
interface MMU : ... bsv:method :: Action addrResponse (Bit ((64) physAddr));
```

```
interface MMU : ... bsv:method :: Action reportStateDbg (DmaDbgRec rec);
```

```
interface MMU : ... bsv:method :: Action reportMemoryTraffic (Bit ((64) words));
```

```
interface MMU : ... bsv:method :: Action error (Bit ((32) code, Bit#(32) sglId, Bit#(64) offset,  
                                         Bit#(64) extra));
```

Pipe Package

The Pipe package is modeled on Bluespec, Inc's PAClib package. It provides functions and modules for composing pipelines of operations.

Pipe Interfaces

```
interface Pipe : PipeIn (type a)
```

Corresponds to the input interface of a FIFO.

```
  enq (a v) → Action
```

```
  notFull → Bool
```

```
interface Pipe : PipeOut (type a)
```

Corresponds to the output interface of a FIFO.

```
  first → a
```

```
  deq → Action
```

```
  notEmpty → Bool
```

```
typeclass Pipe : ToPipeIn (type a, type b)
```

function toPipeIn (*b in*) → PipeIn#(*a*)

Returns a PipeIn to the object “in” with no additional buffering.

typeclass Pipe :: ToPipeOut (*type a, type b*)

function toPipeOut (*b in*) → PipeOut#(*a*)

Returns a PipeOut from the object “in” with no additional buffering.

typeclass Pipe :: MkPipeIn (*type a, type b*)

module mkPipeIn (*b in*) → (PipeIn#(*a*))

Instantiates a module whose interface is a PipeIn to the input parameter “in”. Includes a FIFO buffering stage.

typeclass Pipe :: MkPipeOut (*type a, type b*)

module mkPipeOut (*b in*) → (PipeOut#(*a*))

Instantiates a module whose interface is PipeOut from the input parameter “in”. Includes a FIFO buffering stage.

instance Pipe :: ToPipeIn (*a, FIFO#(a)*)

Converts a FIFO to a PipeIn.

instance Pipe :: ToPipeOut (*a, function a pipefn()*)

Converts a function to a PipeOut.

instance Pipe :: ToPipeOut (*a, Reg#(a)*)

Converts a register to a PipeOut.

instance Pipe :: ToPipeIn (*Vector#(m, a), Gearbox#(m, n, a)*)

Converts a Gearbox to a PipeOut.

instance Pipe :: ToPipeOut (*a, FIFO#(a)*)

Converts a FIFO to a PipeOut.

instance Pipe :: ToPipeOut (*Vector#(n, a), MIMO#(k, n, sz, a)*)

Converts a MIMO to a PipeOut.

instance Pipe :: ToPipeOut (*Vector#(n, a), Gearbox#(m, n, a)*)

Converts a Gearbox to a PipeOut.

instance Pipe :: MkPipeOut (*a, Get#(a)*)

Instantiates a pipelined PipeOut from a Get interface.

instance Pipe :: MkPipeIn (*a, Put#(a)*)

Instantiates a pipelined PipeIn to a Put interface.

Get and Put Pipes

instance Pipe :: ToGet (*PipeOut #(*a*), a*)

instance Pipe :: ToPut (*PipeIn #(*a*), a*)

Connectable Pipes

instance Pipe :: Connectable (*PipeOut#(*a*), Put#(*a*)*)

instance Pipe :: Connectable (*PipeOut#(*a*), PipeIn#(*a*)*)

Mapping over Pipes

function `Pipe::toCountedPipeOut` (*Reg#(Bit#(n)) r, PipeOut#(a) pipe*) → *PipeOut#(a)*

function `Pipe::zipPipeOut` (*PipeOut#(a) ina, PipeOut#(b) inb*) → *PipeOut#(Tuple2#(a,b))*
Returns a *PipeOut* whose elements are 2-tuples of the elements of the input pipes.

function `Pipe::mapPipe` (*function b f(a av), PipeOut#(a) apipe*) → *PipeOut#(b)*
Returns a *PipeOut* that maps the function *f* to each element of the input pipes with no buffering.

module `Pipe::mkMapPipe` (*function b f(a av), PipeOut#(a) apipe*) → *(PipeOut#(b))*
Instantiates a *PipeOut* that maps the function *f* to each element of the input pipes using a FIFO for buffering.

function `Pipe::mapPipeIn` (*function b f(a av), PipeIn#(b) apipe*) → *PipeIn#(a)*
Returns a *PipeIn* applies the function *f* to each value that is enqueued.

Reducing Pipes

Functions on Pipes of Vectors

function `Pipe::unvectorPipeOut` (*PipeOut#(Vector#(l, a)) in*) → *PipeOut#(a)*

Funneling and Unfunneling

module `Pipe::mkFunnel` (*PipeOut#(Vector#(mk, a)) in*) → *(PipeOut#(Vector#(m, a)))*
Returns *k* Vectors of *m* elements for each *Vector#(mk,a)* element of the input pipe.

module `Pipe::mkFunnel1` (*PipeOut#(Vector#(k, a)) in*) → *(PipeOut#(a))*
Same as *mkFunnel*, but returns *k* singleton elements for each vector element of the input pipe.

module `Pipe::mkFunnelGB1` (*Clock slowClock, Reset slowReset, Clock fastClock, Reset fastReset, PipeOut#(Vector#(k, a)) in*) → *(PipeOut#(a))*
Same as *mkFunnel1*, but uses a Gearbox with a 1 to *k* ratio.

module `Pipe::mkUnfunnel` (*PipeOut#(Vector#(m, a)) in*) → *(PipeOut#(Vector#(mk, a)))*
The dual of *mkFunnel*. Consumes *k* elements from the input pipe, each of which is an *m*-element vector, and returns an *mk*-element vector.

module `Pipe::mkUnfunnelGB` (*Clock slowClock, Reset slowReset, Clock fastClock, Reset fastReset, PipeOut#(Vector#(l, a)) in*) → *(PipeOut#(Vector#(k, a)))*
The same as *mkUnfunnel*, but uses a Gearbox with a 1-to-*k*.

module `Pipe::mkRepeat` (*UInt#(n) repetitions, PipeOut#(a) inpipe*) → *(PipeOut#(a))*
Returns a *PipeOut* which repeats each element of the input pipe the specified number of times.

Fork and Join

Fork and Join with limited scalability

module `Pipe::mkForkVector` (*PipeOut#(a) inpipe*) → *(Vector#(n, PipeOut#(a)))*
Replicates each element of the input pipe to each of the output pipes. It uses a FIFO per output pipe.

module `Pipe::mkSizedForkVector` (*Integer size, PipeOut#(a) inpipe*) → *(Vector#(n, PipeOut#(a)))*
Used a *SizedFIFO* for each of the output pipes.

module `Pipe::mkJoin` (*function c f(a av, b bv), PipeOut#(a) apipe, PipeOut#(b) bpipe*) → *(PipeOut#(c))*
Returns a *PipeOut* that applies the function *f* to the elements of the input pipes, with no buffering.

```

module Pipe : mkJoinBuffered (function c f(a av, b bv), PipeOut#(a) apipe, PipeOut#(b) bpipe) →
    (PipeOut#(c))
    Returns a PipeOut that applies the function f to the elements of the input pipes, using a FIFO to buffer the
    output.

```

module Pipe : **mkJoinVector** (*function* $b\ f(\text{Vector}\#(n, a)\ av), \text{Vector}\#(n, \text{PipeOut}\#(a))\ apipes) \rightarrow (\text{PipeOut}\#(b)$
 Same as `mkJoin`, but operates on a vector of `PipeOut` as input.

Funnel Pipes

Fork and Join with tree-based fanout and fanin for scalability.

These are used by MemReadEngine and MemWriteEngine.

```

typedef Pipe::Vector#(j,PipeOut#(a))    FunnelPipe#(numeric type j, numeric type k, type a, nu
typedef Pipe::Vector#(k,PipeOut#(a))    UnFunnelPipe#(numeric type j, numeric type k, type a, nu
typeclass Pipe::FunnelPipesPipelined (numeric type j, numeric type k, type a, numeric type bpc)

```

```

module mkFunnelPipesPipelined (Vector#(k, PipeOut#(a)) in)  $\rightarrow$  (FunnelPipe#(j,k,a,bpc)
module mkFunnelPipesPipelinedRR (Vector#(k, PipeOut#(a)) in, Integer c)  $\rightarrow$  (Fun-
nelPipe#(j,k,a,bpc)
module mkUnFunnelPipesPipelined (Vector#(j, PipeOut#(Tuple2#(Bit#(TLog#(k)), a))) in)  $\rightarrow$ 
(UnFunnelPipe#(j,k,a,bpc)
module mkUnFunnelPipesPipelinedRR (Vector#(j, PipeOut#(a)) in, Integer c)  $\rightarrow$  (UnFun-
nelPipe#(j,k,a,bpc)

```

```
instance Pipe :: FunnelPipesPipelined (l, l, a, bpc)
```

```
instance Pipe :: FunnelPipesPipelined (l, k, a, bpc)
```

```

module Pipe : mkUnFunnelPipesPipelinedInternal (Vector#(1, Pipe-  

                                                    Out#(Tuple2#(Bit#(TLog#(k)), a))) in)  

  → (UnFunnelPipe#(1,k,a,bpc)

```

```
module Pipe :: mkFunnelPipes (Vector#(mk, PipeOut#(a)) ins) → (Vector#(m, PipeOut#(a))
```

```
module Pipe :: mkFunnelPipes1 (Vector#(k, PipeOut#(a)) ins) → (PipeOut#(a)
```

```
module Pipe : mkUnfunnelPipes (Vector#(m, PipeOut#(a)) ins) → (Vector#(mk, PipeOut#(a)))
```

$$\text{module Pipe} :: \text{mkPipelinedForkVector} \text{ (PipeOut\#(a) \quad inpipe, \quad Integer \quad id) \quad \rightarrow \quad (UnFun-}$$

$$\text{nelPipe\#(1,k,a,bpc)}$$

Delimited Pipes

```
interface Pipe :: FirstLastPipe (type a)
```

A pipe whose elements two-tuples of boolean values indicating first and last in a series. The ttype a indicates the type of the counter used.

pipe \rightarrow PipeOut#(Tuple2#(Bool,Bool))
The pipe of delimited elements

start (*a count*) \rightarrow Action
Starts the series of count elements

```

module Pipe : mkFirstLastPipe → (FirstLastPipe#(a))
    Creates a FirstLastPipe.

struct Pipe : RangeConfig#(type a)
    The base, limit and step for mkRangePipeOut.

    xbase → a
    xlimit → a
    xstep → a

interface Pipe : RangePipeIfc (type a)

    pipe → PipeOut#(a)
    isFirst → Bool
    isLast → Bool
    start (RangeConfig#(a) cfg) → Action

module Pipe : mkRangePipeOut → (RangePipeIfc#(a))
    Creates a Pipe of values from xbase to xlimit by xstep. Used by MemRead.

```

Portal Package

PipePortal Interface

```

interface Portal : PipePortal (numeric type numRequests, numeric type numIndications, numeric type
    slaveDataWidth)

    messageSize (Bit#(16) methodNumber) → Bit#(16)
        Returns the message size of the methodNumber method of the portal.

    requests → Vector#(numRequests, PipeIn#(Bit#(slaveDataWidth)))
    indications → Vector#(numIndications, PipeOut#(Bit#(slaveDataWidth)))

```

MemPortal Interface

```

interface Portal : MemPortal (numeric type slaveAddrWidth, numeric type slaveDataWidth)

    slave → PhysMemSlave#(slaveAddrWidth,slaveDataWidth)
    interrupt → ReadOnly#(Bool)
    top → WriteOnly#(Bool)

function Portal : getSlave (MemPortal#(_a, _d) p) → PhysMemSlave(_a, _d)
function Portal : getInterrupt (MemPortal#(_a, _d) p) → ReadOnly#(Bool)
function Portal : getInterruptVector (Vector#(numPortals, MemPortal#(_a, _d)) portals) → Vec-
    tor#(16,ReadOnly#(Bool))

```

ShareMemoryPortal Interface

interface Portal :: **SharedMemoryPortal** (numeric type dataBusWidth)

Should be in SharedMemoryPortal.bsv

readClient → MemReadClient(dataBusWidth)

writeClient → MemWriteClient#(dataBusWidth)

cfg → SharedMemoryPortalConfig

interrupt → ReadOnly#(Bool)

ConnectalTop Interface

interface Portal :: **ConnectalTop** (numeric type addrWidth, numeric type dataWidth, type pins, numeric type numMasters)

Interface ConnectalTop is the interface exposed by the top module of a Connectal hardware design.

slave → PhysMemSlave#(32,32)

masters → Vector#(numMasters,PhysMemMaster#(addrWidth, dataWidth))

interrupt → Vector#(16,ReadOnly#(Bool))

leds → LEDS

pins → pins

StdConnectalTop Typedef

typedef Portal :: **StdConnectalTop** (numeric type addrWidth) → Connectal-
Top#(addrWidth,64,Empty,0)

Type StdConnectalTop indicates a Connectal hardware design with no user defined pins and no user of host shared memory. The “pins” interface is Empty and the number of masters is 0.

typedef Portal :: **StdConnectalDmaTop** (numeric type addrWidth) → Connectal-
Top#(addrWidth,64,Empty,1)

Type StdConnectalDmaTop indicates a Connectal hardware design with no user defined pins and a single client of host shared memory. The “pins” interface is Empty and the number of masters is 1.

Connectal C/C++ Libraries

C/C++ Portal

Connecting to Bluesim

BLUESIM_SOCKET_NAME

Controls the name of the socket used for connecting software and hardware simulated by bluesim.

Connecting to Xsim and Verilator

SOFTWARE_SOCKET_NAME

Controls the name of the socket used for connecting software and hardware simulated by xsim/verilator.

Automatically Programming the FPGA

Connectal application executables or shared objects contain the FPGA bitstream in the “fpgadata” section of the ELF file. When the application (or library) first tries to access the hardware, the Connectal library automatically programs the FPGA with the associated bitstream, unless *noprogram* is set to a non-zero value or environment variable *NOPROGRAM* is nonzero.

In the case of simulation hardware, the simulator is launched when the application first tries to access the hardware. This behavior is also suppressed by a nonzero value for either *noprogram* or *NOPROGRAM*.

int **noprogram**

If *noprogram* is set to a non-zero value, then the FPGA is not programmed automatically.

Tracing Simulation

DUMP_VCD

If set, directs the simulator to dump a VCD trace to the \$DUMP_VCD.

int **simulator_dump_vcd**

The application can enable VCD tracing by setting *simulator_dump_vcd* to 1. It takes the file name from *simulator_vcd_name*. **DUMP_VCD** overrides this variable.

const char ***simulator_vcd_name;**

Specifies the name of the vcd file. Defaults to “dump.vcd”. **DUMP_VCD** overrides this variable.

Zynq Clock Control

void **setClockFrequency** (int *clkNum*, long *requestedFrequency*, long **actualFrequency*)

Changes the frequency of Zynq FPGA Clock *clkNum* to the closest frequency to *requestedFrequency* available from the PLL. If the *actualFrequency* pointer is non-null, stores the actual frequency before returning.

Portal Memory

int **portalAlloc** (size_t *size*, int *cached*)

Uses portalmem to allocate a region of *size* bytes.

On platforms that support non-cache-coherent I/O (e.g., zedboard), *cached*=0 indicates that the programmable logic will use a port to memory that is not snooped by the CPU’s caches. In this case, it is up to the allocation to flush or invalidate the CPU cache as needed, using **portalCacheFlush**().

Returns the file descriptor associated with the memory region.

void ***portalMmap** (int *fd*, size_t *size*)

Memory maps *size* bytes of the portal memory region indicated by *fd*.

Returns a pointer to memory on success or -1 on failure.

portalCacheFlush (int *fd*, void **__p*, long *size*, int *flush*)

PortalPoller

class **PortalPoller**

Polls portals

PortalPoller::PortalPoller(int *autostart* = 1)

If *autostart* is 1, then invoke *start* () from *registerInstance* ()

void PortalPoller::start();

Starts the poller. Called automatically from `registerInstance()` if `autostart` is 1.

void PortalPoller::stop();

Stops the poller.

int PortalPoller::timeout

The timeout value, in milliseconds, passed to **`:c:function:'poll()'`**

Deprecated Functions

void *portalExec (void *__x)

Polls the registered portals and invokes their callback **handlers**. ()

void portalExec_start ()

void portalExec_poll ()

Connectal Examples

Simple Example

Indices and tables

- [genindex](#)
- [modindex](#)
- [search](#)

a

AddressGenerator, 45
Arith, 45
AxiStream, 46

c

CtrlMux, 47

h

HostInterface, 48

l

Leds, 48

m

MemPortal, 48
MemReadEngine, 49
MemTypes, 49
MMU, 53

p

Pipe, 55
Portal, 59

Symbols

-O os, -option=os
 command line option, 39
 -bsvpath=dirname
 command line option, 40
 -cachedir=dirname
 command line option, 39
 -cflags=flags
 command line option, 39
 -constraint=file.xdc
 command line option, 40
 -derivedclockperiod
 command line option, 41
 -dump_map
 command line option, 39
 -implconstraint=file.xdc
 command line option, 40
 -interfaces=...
 command line option, 39
 -ipdir=dir
 command line option, 39
 -mainclockperiod
 command line option, 41
 -nocache
 command line option, 39
 -nohardware
 command line option, 40
 -nonstrict
 command line option, 39
 -pinfo=project.json
 command line option, 40
 -pinout=pinusage.json
 command line option, 40
 -project-dir=dir
 command line option, 39
 -prtop=mkTop.dcp
 command line option, 41
 -prvariant=name
 command line option, 41
 -qip=core.qip
 command line option, 41
 -qsf=settings.qsf

 command line option, 41
 -reconfig=modulenames
 command line option, 41
 -shared
 command line option, 40
 -stl=stltype
 command line option, 40
 -xci=core.xci
 command line option, 40
 -xelabflags=flags
 command line option, 41
 -xsimflags
 command line option, 41
 -B board, -board=board
 command line option, 39
 -D var=value
 command line option, 39
 -I dir -cinclude=dir
 command line option, 39
 -L libdir, -clibdir=libdir
 command line option, 39
 -P modulename, -partition=modulename
 command line option, 40
 -S libfile, -clibfiles=libfile.a
 command line option, 39
 -V file.v, -verilog=file.v
 command line option, 40
 -b options, -bscflags=options
 command line option, 40
 -l lib, -clib=lib
 command line option, 39
 -m foo.cpp, -bsimsource foo.cpp
 command line option, 41
 -q, -qtused
 command line option, 41
 -s foo.cpp, -source=foo.cpp
 command line option, 39
 -v, -verbose
 command line option, 39
 .. bsv:function:: atype toAxiStream(btype b) (typeclass in
 package AxiStream), 47
 .. bsv:method:: Action addrResponse(Bit (interface in

package MMU), 55
 .. bsv:method:: Action error(Bit (interface in package MMU), 55
 .. bsv:method:: Action idResponse(Bit (interface in package MMU), 54
 .. bsv:method:: Action reportMemoryTraffic(Bit (interface in package MMU), 55
 .. bsv:method:: Action reportStateDbg(DmaDbgRec rec) (interface in package MMU), 55
 12 SGListPageShift0 (typedef in package MMU), 53
 16 SGListPageShift4 (typedef in package MMU), 53
 20 SGListPageShift8 (typedef in package MMU), 53
 24 SGListPageShift12 (typedef in package MMU), 53
 32 MaxNumSGLists (typedef in package MMU), 53
 32 MemServerTags (typedef in package MemTypes), 49
 44 MemOffsetSize (typedef in package MemTypes), 49
 6 MemTagSize (typedef in package MemTypes), 49
 8 BurstLenSize (typedef in package MemTypes), 49
 8 IndexWidth (typedef in package MMU), 53

tready() (AxiStream::AxiStreamMaster method), 46
 tvalid() (AxiStream::AxiStreamMaster method), 46
 tlast() (AxiStream::AxiStreamMaster method), 46
 tkeep() (AxiStream::AxiStreamMaster method), 46

A

a) vadd (function in package Arith), 46
 a) vmul (function in package Arith), 46
 add (function in package Arith), 45
 AddrBeat#(numeric type addrWidth) (struct in package AddressGenerator), 45
 AddressGenerator (interface in package AddressGenerator), 45
 AddressGenerator (package), 45
 addrTrans() (MMU::MemServerRequest method), 55
 AddrTransRequest (struct in package MMU), 53
 android.exe
 Make Target, 36, 37
 Arith (package), 45
 AxiStream (package), 46
 AxiStreamMaster (interface in package AxiStream), 46
 AxiStreamSlave (interface in package AxiStream), 46

B

Bit# (function in package Arith), 46
 Bit#(32) SGLId (typedef in package MemTypes), 49
 Bit#(TLog#(MaxNumSGLists)) RegionsIdx (typedef in package MMU), 53
 Bit#(TLog#(MaxNumSGLists)) SGLId (typedef in package MMU), 53
 bits
 Make Target, 36, 37
 BOARD

Make Variable, 34
 booland (function in package Arith), 45
 boolor (function in package Arith), 45
 BSC_CXXFLAGS
 Make Variable, 38
 BSCFLAGS_COMMON
 Make Variable, 38
 BSCFLAGS_EXTRA
 Make Variable, 38
 BSCFLAGS_PROJECT
 Make Variable, 38
 BSCOPTFLAGS
 Make Variable, 38
 BSCPATHFLAGS
 Make Variable, 38
 bsim
 Make Target, 36, 37
 bsim_exe
 Make Target, 36, 37
 BsimHost (interface in package HostInterface), 48
 bsvfile
 command line option, 39
 BSVFILES
 Make Variable, 35
 build.%
 Make Target, 36
 BUILDDCACHE
 Make Variable, 35
 BUILDDCACHE_CACHEDIR
 Make Variable, 35
 BURST_LEN_SIZE
 Make Variable, 36

C

command line option
 -O os, -option=os, 39
 -bsvpath=dirname, 40
 -cachedir=dirname, 39
 -cflags=flags, 39
 -constraint=file.xdc, 40
 -derivedclockperiod, 41
 -dump_map, 39
 -implconstraint=file.xdc, 40
 -interfaces=..., 39
 -ipdir=dir, 39
 -mainclockperiod, 41
 -nocache, 39
 -nohardware, 40
 -nonstrict, 39
 -pinfo=project.json, 40
 -pinout=pinusage.json, 40
 -project-dir=dir, 39
 -prtop=mkTop.dcp, 41
 -prvariant=name, 41

- qip=core.qip, 41
- qsf=settings.qsf, 41
- reconfig=modulenames, 41
- shared, 40
- stl=stltype, 40
- xci=core.xci, 40
- xelabflags=flags, 41
- xsimflags, 41
- B board, -board=board, 39
- D var=value, 39
- I dir -cinclude=dir, 39
- L libdir, -clibdir=libdir, 39
- P modulename, -partition=modulename, 40
- S libfile, -clibfiles=libfile.a, 39
- V file.v, -verilog=file.v, 40
- b options, -bscflags=options, 40
- l lib, -clib=lib, 39
- m foo.cpp, -bsimsource foo.cpp, 41
- q, -qtused, 41
- s foo.cpp, -source=foo.cpp, 39
- v, -verbose, 39
- bsvfile, 39
- configResp() (MMU::MMUIndication method), 54
- Connectable (instance in package AxiStream), 47
- Connectable (instance in package MemTypes), 53
- Connectable (instance in package Pipe), 56
- CONNECTAL_DEBUG
 - Make Variable, 38
- CONNECTAL_NDK_PARAM
 - Make Variable, 37
- CONNECTAL_NOHARDWARE
 - Make Variable, 37
- CONNECTAL_SHARED
 - Make Variable, 37
- CONNECTALFLAGS
 - Make Variable, 35
- ConnectalTop (interface in package Portal), 60
- CPPFILES
 - Make Variable, 35
- CPPFILES2
 - Make Variable, 35
- CtrlMux (package), 47

D

- DataBusWidth (typedef in package HostInterface), 48
- dbg() (MemTypes::DmaDbg method), 53
- DEFAULT_DERIVED_CLOCK_PERIOD
 - Make Variable, 35
- deq() (Pipe::PipeOut method), 55
- DERIVED_CLOCK_PERIOD
 - Make Variable, 36
- DmaDbg (interface in package MemTypes), 53
- DmaErrorType (struct in package MMU), 54

E

- enq() (Pipe::PipeIn method), 55
- environment variable
 - BLUESIM_SOCKET_NAME, 60
 - CONNECTALDIR, 34
 - DUMP_VCD, 61
 - NOFFGAJTAG, 37
 - NOPROGRAM, 37, 61
 - RUNPARAM, 37
 - SOFTWARE_SOCKET_NAME, 60
- eq (function in package Arith), 45
- error() (MMU::MMUIndication method), 54
- exe
 - Make Target, 37

F

- first() (Pipe::PipeOut method), 55
- FirstLastPipe (interface in package Pipe), 58
- FunnelPipesPipelined (instance in package Pipe), 58
- FunnelPipesPipelined (typeclass in package Pipe), 58
- FunnelPipesPipelined.mkFunnelPipesPipelined (module in package Pipe), 58
- FunnelPipesPipelined.mkFunnelPipesPipelinedRR (module in package Pipe), 58
- FunnelPipesPipelined.mkUnFunnelPipesPipelined (module in package Pipe), 58
- FunnelPipesPipelined.mkUnFunnelPipesPipelinedRR (module in package Pipe), 58

G

- gentarget
 - Make Target, 36
- getInterrupt (function in package Portal), 59
- getInterruptVector (function in package Portal), 59
- getMemoryTraffic() (MemTypes::DmaDbg method), 53
- getSlave (function in package Portal), 59

H

- H2S_INTERFACES
 - Make Variable, 35
- handlers. (C function), 62
- HostInterface (package), 48

I

- idRequest() (MMU::MMURequest method), 54
- idReturn() (MMU::MMURequest method), 54
- INTERFACES
 - Make Variable, 34
- IPDIR
 - Make Variable, 35

isFirst() (Pipe::RangePipeIfc method), 59

isLast() (Pipe::RangePipeIfc method), 59

L

LEDS (interface in package Leds), 48

Leds (package), 48

leds() (in package Leds), 48

LedsWidth (typedef in package Leds), 48

M

MAIN_CLOCK_PERIOD

Make Variable, 35

Make Target

android.exe, 36, 37

bits, 36, 37

bsim, 36, 37

bsim_exe, 36, 37

build.%, 36

exe, 37

gentarget, 36

prebuild, 36

run, 37

run.%, 36

ubuntu.exe, 36, 37

verilog, 36

xsim, 36, 37

Make Variable

BOARD, 34

BSC_CXXFLAGS, 38

BSCFLAGS_COMMON, 38

BSCFLAGS_EXTRA, 38

BSCFLAGS_PROJECT, 38

BSCOPTFLAGS, 38

BSCPATHFLAGS, 38

BSVFILES, 35

BUILDCACHE, 35

BUILDCACHE_CACHEDIR, 35

BURST_LEN_SIZE, 36

CONNECTAL_DEBUG, 38

CONNECTAL_NDK_PARAM, 37

CONNECTAL_NOHARDWARE, 37

CONNECTAL_SHARED, 37

CONNECTALFLAGS, 35

CPPFILES, 35

CPPFILES2, 35

DEFAULT_DERIVED_CLOCK_PERIOD, 35

DERIVED_CLOCK_PERIOD, 36

H2S_INTERFACES, 35

INTERFACES, 34

IPDIR, 35

MAIN_CLOCK_PERIOD, 35

MEM_READ_INTERFACES, 35

MEM_WRITE_INTERFACES, 35

NUMBER_OF_MASTERS, 34

PIN_TYPE, 34

PIN_TYPE_INCLUDE, 34

PINOUT_FILE, 35

PORTAL_DUMP_MAP, 35

RUN_ARGS, 38

RUNPARAM, 36

S2H_INTERFACES, 35

USE_BUILDCACHE, 35, 38

V, 35, 37

VIVADOFLAGS, 38

mapPipe (function in package Pipe), 57

mapPipeIn (function in package Pipe), 57

MEM_READ_INTERFACES

Make Variable, 35

MEM_WRITE_INTERFACES

Make Variable, 35

MemData#(numeric type dsz) (struct in package MemTypes), 50

MemDataF#(numeric type dsz) (struct in package MemTypes), 50

MemengineCmd (struct in package MemTypes), 52

memoryTraffic() (MMU::MemServerRequest method), 55

MemPortal (interface in package Portal), 59

MemPortal (package), 48

MemReadClient (interface in package MemTypes), 51

MemReadEngine (interface in package MemTypes), 52

MemReadEngine (package), 49

MemReadEngineServer (interface in package MemTypes), 52

MemReadServer (interface in package MemTypes), 51

MemRequest (struct in package MemTypes), 49

MemServerIndication

(interface in package MMU), 55

MemServerRequest

(interface in package MMU), 55

MemTypes (package), 49

MemWriteClient (interface in package MemTypes), 51

MemWriteEngine (interface in package MemTypes), 52

MemWriteEngineServer (interface in package MemTypes), 52

MemWriteServer (interface in package MemTypes), 51

messageSize() (Portal::PipePortal method), 59

mkAddressGenerator (module in package AddressGenerator), 45

MkAxiStream (instance in package AxiStream), 47

MkAxiStream (typeclass in package AxiStream), 47

MkAxiStream.mkAxiStream (module in package AxiStream), 47

mkFirstLastPipe (module in package Pipe), 58

mkForkVector (module in package Pipe), 57

mkFunnel (module in package Pipe), 57

mkFunnel1 (module in package Pipe), 57

mkFunnelGB1 (module in package Pipe), 57

mkFunnelPipes (module in package Pipe), 58
 mkFunnelPipes1 (module in package Pipe), 58
 mkInterruptMux (module in package CtrlMux), 47
 mkJoin (module in package Pipe), 57
 mkJoinBuffered (module in package Pipe), 57
 mkJoinVector (module in package Pipe), 58
 mkMapPipe (module in package Pipe), 57
 mkMemPortal (module in package MemPortal), 48
 mkMemReadEngine(MemReadEngine (module in package MemReadEngine), 49
 mkMemReadEngineBuff (module in package MemReadEngine), 49
 mkMMU (module in package MMU), 55
 MkPipeIn (instance in package Pipe), 56
 MkPipeIn (typeclass in package Pipe), 56
 MkPipeIn.mkPipeIn (module in package Pipe), 56
 mkPipelinedForkVector (module in package Pipe), 58
 MkPipeOut (instance in package Pipe), 56
 MkPipeOut (typeclass in package Pipe), 56
 MkPipeOut.mkPipeOut (module in package Pipe), 56
 mkRangePipeOut (module in package Pipe), 59
 mkRepeat (module in package Pipe), 57
 mkSizedForkVector (module in package Pipe), 57
 mkSlaveMux (module in package CtrlMux), 47
 mkUnfunnel (module in package Pipe), 57
 mkUnfunnelGB (module in package Pipe), 57
 mkUnfunnelPipes (module in package Pipe), 58
 mkUnFunnelPipesPipelinedInternal (module in package Pipe), 58
 MMU (interface in package MMU), 53
 MMU (package), 53
 MMUIndication
 (interface in package MMU), 54
 MMURequest
 (interface in package MMU), 53
 mul (function in package Arith), 45

N

NOPROGRAM, 61
 noprogram (C variable), 61
 notEmpty() (Pipe::PipeOut method), 55
 notFull() (Pipe::PipeIn method), 55
 NUMBER_OF_MASTERS
 Make Variable, 34
 NumberOfMasters (typedef in package HostInterface), 48

P

PcieHost (interface in package HostInterface), 48
 PhysAddrWidth (typedef in package HostInterface), 48
 PhysMemMaster (interface in package MemTypes), 50
 PhysMemReadClient (interface in package MemTypes), 50
 PhysMemReadServer (interface in package MemTypes), 51

PhysMemRequest#(numeric type addrWidth, dataWidth) (struct in package MemTypes), 49
 PhysMemSlave (interface in package MemTypes), 51
 PhysMemWriteClient (interface in package MemTypes), 51
 PhysMemWriteServer (interface in package MemTypes), 51
 PIN_TYPE
 Make Variable, 34
 PIN_TYPE_INCLUDE
 Make Variable, 34
 PINOUT_FILE
 Make Variable, 35
 PinType (typedef in package HostInterface), 48
 Pipe (package), 55
 PipeIn (interface in package Pipe), 55
 PipeOut (interface in package Pipe), 55
 PipePortal (interface in package Portal), 59
 Portal (package), 59
 PORTAL_DUMP_MAP
 Make Variable, 35
 portalAlloc (C function), 61
 portalCacheFlush (C function), 61
 portalExec (C function), 62
 portalExec_poll (C function), 62
 portalExec_start (C function), 62
 portalMmap (C function), 61
 PortalPoller (C++ class), 61
 PortalPoller::PortalPoller::timeout (C++ member), 62
 prebuild
 Make Target, 36

R

RangeConfig#(type a) (struct in package Pipe), 59
 RangePipeIfc (interface in package Pipe), 59
 region() (MMU::MMURequest method), 54
 rshift (function in package Arith), 46
 run
 Make Target, 37
 run.%
 Make Target, 36
 RUN_ARGS
 Make Variable, 38
 RUNPARAM
 Make Variable, 36

S

S2H_INTERFACES
 Make Variable, 35
 setClockFrequency (C function), 61
 setInterface() (MMU::MMURequest method), 54
 setTileState() (MMU::MemServerRequest

method), 55
sglist() (MMU::MMURequest
method), 54
SharedMemoryPortal (interface in package Portal), 60
simulator_dump_vcd (C variable), 61
start() (Pipe::FirstLastPipe method), 58
start() (Pipe::RangePipeIfc method), 59
stateDbg() (MMU::MemServerRequest
method), 55
StdConnectalDmaTop (typedef in package Portal), 60
StdConnectalTop (typedef in package Portal), 60

T

tdata() (AxiStream::AxiStreamMaster method), 46
tdata() (AxiStream::AxiStreamSlave method), 46
TDiv#(DataBusSize,8) ByteEnableSize (typedef in package MemTypes), 49
tkeep() (AxiStream::AxiStreamSlave method), 46
tlast() (AxiStream::AxiStreamSlave method), 46
ToAxiStream (typeclass in package AxiStream), 47
toCountedPipeOut (function in package Pipe), 57
ToGet (instance in package Pipe), 56
ToGetM (instance in package AxiStream), 47
ToPipeIn (instance in package Pipe), 56
ToPipeIn (typeclass in package Pipe), 55
ToPipeIn.toPipeIn (function in package Pipe), 55
ToPipeOut (instance in package Pipe), 56
ToPipeOut (typeclass in package Pipe), 56
ToPipeOut.toPipeOut (function in package Pipe), 56
ToPut (instance in package Pipe), 56
ToPutM (instance in package AxiStream), 47
treedy() (AxiStream::AxiStreamSlave method), 46
tvalid() (AxiStream::AxiStreamSlave method), 46

U

ubuntu.exe
Make Target, 36, 37
unvectorPipeOut (function in package Pipe), 57
USE_BUILDDCACHE
Make Variable, 35, 38

V

V

Make Variable, 35, 37
Vector#(j,PipeOut#(a)) FunnelPipe#(numeric type j, numeric type k, type a, numeric type bitsPerCycle) (typedef in package Pipe), 58
Vector#(k,PipeOut#(a)) UnFunnelPipe#(numeric type j, numeric type k, type a, numeric type bitsPerCycle) (typedef in package Pipe), 58
verilog
Make Target, 36
VIVADOFLAGS
Make Variable, 38

X

xsim
Make Target, 36, 37

Z

zipPipeOut (function in package Pipe), 57
ZynqHost (interface in package HostInterface), 48